

Google Data Analytics Certificate Capstone (Bellabeat Case Study)

Gabriel Fernandez

```
# Set default options for code chunks
knitr::opts_chunk$set(
  echo = TRUE,          # Display R code and its output
  comment=NA,           # Suppress code comments in output
  warning = FALSE,      # Suppress warning messages
  fig.align='center',   # Align figures in the center
  eval = TRUE           # Evaluate R code
)
```

Prepare and preprocess phase

```
# Import libraries
library(tidyverse) # includes ggplot2
library(skimr) # provides a compact and informative summary of your dataframe or dataset
library(lubridate)
library(janitor) # set of utility functions for data cleaning and data frame tidying tasks
library(RColorBrewer) # Color palettes for data visualization
library(ggcorrplot) # Visualize correlation matrices using ggplot2
library(scales) # formatting and transforming data for visualizations

# display.brewer.all(colorblindFriendly = TRUE)
```

Load datasets

- [Datasets:](#)

These datasets originate from a survey distributed on Amazon Mechanical Turk from 03.12.2016 to 05.12.2016. They include personal tracker data from 30 Fitbit users, covering physical activity, heart rate, and sleep monitoring, with differentiation based on Fitbit types and user behavior.

- Metadata: [Fitbit data dictionary](#)

```
# Clean environment
rm(list = ls())

daily_activity <-
  read_csv("original_data/dailyActivity_merged.csv",
    trim_ws = TRUE,
    show_col_types = FALSE
  )

daily_sleep <- read_csv("original_data/sleepDay_merged.csv",
  trim_ws = TRUE,
```

```

    show_col_types = FALSE
  )

hourly_calories <-
  read_csv("original_data/hourlyCalories_merged.csv",
    trim_ws = TRUE,
    show_col_types = FALSE
  )

hourly_intensities <-
  read_csv("original_data/hourlyIntensities_merged.csv",
    trim_ws = TRUE,
    show_col_types = FALSE
  )

hourly_steps <-
  read_csv("original_data/hourlySteps_merged.csv",
    trim_ws = TRUE,
    show_col_types = FALSE
  )

minute_sleep <-
  read_csv("original_data/minuteSleep_merged.csv",
    trim_ws = TRUE,
    show_col_types = FALSE
  )

seconds_hearttrate <-
  read_csv("original_data/hearttrate_seconds_merged.csv",
    trim_ws = TRUE,
    show_col_types = FALSE
  )

weight_logs <-
  read_csv("original_data/weightLogInfo_merged.csv",
    trim_ws = TRUE,
    show_col_types = FALSE
  )

# Remove trailing spaces (trim_ws = TRUE)

```

dataset

Clean datasets

Clean the daily__activity dataset

```

# Check daily_activity dataset before cleaning
glimpse(daily_activity)

```

```

Rows: 940
Columns: 15
$ Id                <dbl> 1503960366, 1503960366, 1503960366, 150396036~
$ ActivityDate      <chr> "4/12/2016", "4/13/2016", "4/14/2016", "4/15/~
$ TotalSteps        <dbl> 13162, 10735, 10460, 9762, 12669, 9705, 13019~

```

```

$ TotalDistance      <dbl> 8.50, 6.97, 6.74, 6.28, 8.16, 6.48, 8.59, 9.8~
$ TrackerDistance    <dbl> 8.50, 6.97, 6.74, 6.28, 8.16, 6.48, 8.59, 9.8~
$ LoggedActivitiesDistance <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ~
$ VeryActiveDistance <dbl> 1.88, 1.57, 2.44, 2.14, 2.71, 3.19, 3.25, 3.5~
$ ModeratelyActiveDistance <dbl> 0.55, 0.69, 0.40, 1.26, 0.41, 0.78, 0.64, 1.3~
$ LightActiveDistance <dbl> 6.06, 4.71, 3.91, 2.83, 5.04, 2.51, 4.71, 5.0~
$ SedentaryActiveDistance <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ~
$ VeryActiveMinutes   <dbl> 25, 21, 30, 29, 36, 38, 42, 50, 28, 19, 66, 4~
$ FairlyActiveMinutes <dbl> 13, 19, 11, 34, 10, 20, 16, 31, 12, 8, 27, 21~
$ LightlyActiveMinutes <dbl> 328, 217, 181, 209, 221, 164, 233, 264, 205, ~
$ SedentaryMinutes    <dbl> 728, 776, 1218, 726, 773, 539, 1149, 775, 818~
$ Calories             <dbl> 1985, 1797, 1776, 1745, 1863, 1728, 1921, 203~

```

```
# Check missing values and duplicates
```

```

cat(
  "\n",
  "Missing values:",
  sum(is.na(daily_activity)),
  "\n",
  "Duplicate values:",
  sum(duplicated(daily_activity)),
  "\n",
  "Unique Ids:",
  n_distinct(daily_activity$Id)
)

```

Missing values: 0

Duplicate values: 0

Unique Ids: 33

Let us clean:

- Change column names to lower case because R is case sensitive.
- Change “Id” from double to a character because the number represents a category.
- Change “ActivityDate” from char to date.

```
# Clean daily_activity dataset
```

```

daily_activity <-
  # Clean column names
  clean_names(daily_activity) %>%
  # Correct column types
  mutate(id = as.character(id)) %>% # from double to chr
  mutate(activity_date = as.Date(activity_date,
                                format = "%m/%d/%Y")) %>% # from chr to date

  # Remove duplicate rows
  distinct()

# Check daily_activity dataset after cleaning
glimpse(daily_activity)

```

Rows: 940

Columns: 15

```

$ id      <chr> "1503960366", "1503960366", "1503960366", "~
$ activity_date <date> 2016-04-12, 2016-04-13, 2016-04-14, 2016-0~

```

```

$ total_steps          <dbl> 13162, 10735, 10460, 9762, 12669, 9705, 130~
$ total_distance       <dbl> 8.50, 6.97, 6.74, 6.28, 8.16, 6.48, 8.59, 9~
$ tracker_distance     <dbl> 8.50, 6.97, 6.74, 6.28, 8.16, 6.48, 8.59, 9~
$ logged_activities_distance <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
$ very_active_distance <dbl> 1.88, 1.57, 2.44, 2.14, 2.71, 3.19, 3.25, 3~
$ moderately_active_distance <dbl> 0.55, 0.69, 0.40, 1.26, 0.41, 0.78, 0.64, 1~
$ light_active_distance <dbl> 6.06, 4.71, 3.91, 2.83, 5.04, 2.51, 4.71, 5~
$ sedentary_active_distance <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
$ very_active_minutes   <dbl> 25, 21, 30, 29, 36, 38, 42, 50, 28, 19, 66,~
$ fairly_active_minutes <dbl> 13, 19, 11, 34, 10, 20, 16, 31, 12, 8, 27, ~
$ lightly_active_minutes <dbl> 328, 217, 181, 209, 221, 164, 233, 264, 205~
$ sedentary_minutes     <dbl> 728, 776, 1218, 726, 773, 539, 1149, 775, 8~
$ calories              <dbl> 1985, 1797, 1776, 1745, 1863, 1728, 1921, 2~

```

Check missing values and duplicates after cleaning

```

cat("\n",
  "Missing values:",
  sum(is.na(daily_activity)),
  "\n",
  "Duplicate values:",
  sum(duplicated(daily_activity)))

```

Missing values: 0

Duplicate values: 0

Let us print summary statistics to have a better idea of the dataset

```

daily_activity %>%
  summary()

```

```

      id      activity_date      total_steps      total_distance
Length:940   Min.   :2016-04-12   Min.   :    0   Min.   : 0.000
Class :character 1st Qu.:2016-04-19 1st Qu.: 3790 1st Qu.: 2.620
Mode  :character Median :2016-04-26 Median : 7406 Median : 5.245
      Mean   :2016-04-26 Mean   : 7638 Mean   : 5.490
      3rd Qu.:2016-05-04 3rd Qu.:10727 3rd Qu.: 7.713
      Max.   :2016-05-12 Max.   :36019 Max.   :28.030

tracker_distance logged_activities_distance very_active_distance
Min.   : 0.000   Min.   :0.0000   Min.   : 0.000
1st Qu.: 2.620   1st Qu.:0.0000   1st Qu.: 0.000
Median : 5.245   Median :0.0000   Median : 0.210
Mean   : 5.475   Mean   :0.1082   Mean   : 1.503
3rd Qu.: 7.710   3rd Qu.:0.0000   3rd Qu.: 2.053
Max.   :28.030   Max.   :4.9421   Max.   :21.920

moderately_active_distance light_active_distance sedentary_active_distance
Min.   :0.0000   Min.   : 0.000   Min.   :0.000000
1st Qu.:0.0000   1st Qu.: 1.945   1st Qu.:0.000000
Median :0.2400   Median : 3.365   Median :0.000000
Mean   :0.5675   Mean   : 3.341   Mean   :0.001606
3rd Qu.:0.8000   3rd Qu.: 4.782   3rd Qu.:0.000000
Max.   :6.4800   Max.   :10.710   Max.   :0.110000

very_active_minutes fairly_active_minutes lightly_active_minutes
Min.   : 0.00   Min.   : 0.00   Min.   : 0.0
1st Qu.: 0.00   1st Qu.: 0.00   1st Qu.:127.0
Median : 4.00   Median : 6.00   Median :199.0
Mean   : 21.16   Mean   :13.56   Mean   :192.8

```

3rd Qu.: 32.00	3rd Qu.: 19.00	3rd Qu.:264.0
Max. :210.00	Max. :143.00	Max. :518.0
sedentary_minutes	calories	
Min. : 0.0	Min. : 0	
1st Qu.: 729.8	1st Qu.:1828	
Median :1057.5	Median :2134	
Mean : 991.2	Mean :2304	
3rd Qu.:1229.5	3rd Qu.:2793	
Max. :1440.0	Max. :4900	

- This summary helps us explore each attribute quickly. We notice that some attributes have a minimum value of zero (total_step, total_distance, calories). Let us explore this observation.

```
# Check where total_steps is zero
filter(daily_activity, total_steps == 0)
```

```
# A tibble: 77 x 15
  id      activity_~1 total~2 total~3 track~4 logge~5 very_~6 moder~7 light~8
  <chr>    <date>         <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>
1 1503960366 2016-05-12         0      0      0      0      0      0      0
2 1844505072 2016-04-24         0      0      0      0      0      0      0
3 1844505072 2016-04-25         0      0      0      0      0      0      0
4 1844505072 2016-04-26         0      0      0      0      0      0      0
5 1844505072 2016-05-02         0      0      0      0      0      0      0
6 1844505072 2016-05-07         0      0      0      0      0      0      0
7 1844505072 2016-05-08         0      0      0      0      0      0      0
8 1844505072 2016-05-09         0      0      0      0      0      0      0
9 1844505072 2016-05-10         0      0      0      0      0      0      0
10 1844505072 2016-05-11         0      0      0      0      0      0      0
# ... with 67 more rows, 6 more variables: sedentary_active_distance <dbl>,
# very_active_minutes <dbl>, fairly_active_minutes <dbl>,
# lightly_active_minutes <dbl>, sedentary_minutes <dbl>, calories <dbl>, and
# abbreviated variable names 1: activity_date, 2: total_steps,
# 3: total_distance, 4: tracker_distance, 5: logged_activities_distance,
# 6: very_active_distance, 7: moderately_active_distance,
# 8: light_active_distance
```

- We found 77 observations where total_step is zero. We should delete these observations so they do not affect our mean and median. If the total_step is zero, the person did not wear the Fitbit.

```
# Check where calories is zero
filter(daily_activity, calories == 0)
```

```
# A tibble: 4 x 15
  id      activity_~1 total~2 total~3 track~4 logge~5 very_~6 moder~7 light~8
  <chr>    <date>         <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>
1 1503960366 2016-05-12         0      0      0      0      0      0      0
2 6290855005 2016-05-10         0      0      0      0      0      0      0
3 8253242879 2016-04-30         0      0      0      0      0      0      0
4 8583815059 2016-05-12         0      0      0      0      0      0      0
# ... with 6 more variables: sedentary_active_distance <dbl>,
# very_active_minutes <dbl>, fairly_active_minutes <dbl>,
# lightly_active_minutes <dbl>, sedentary_minutes <dbl>, calories <dbl>, and
# abbreviated variable names 1: activity_date, 2: total_steps,
# 3: total_distance, 4: tracker_distance, 5: logged_activities_distance,
# 6: very_active_distance, 7: moderately_active_distance,
# 8: light_active_distance
```

```
# Check where total_distance is zero
filter(daily_activity, total_distance == 0)
```

```
# A tibble: 78 x 15
```

	id	activity~1	total~2	total~3	track~4	logge~5	very_~6	moder~7	light~8
	<chr>	<date>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	1503960366	2016-05-12	0	0	0	0	0	0	0
2	1844505072	2016-04-24	0	0	0	0	0	0	0
3	1844505072	2016-04-25	0	0	0	0	0	0	0
4	1844505072	2016-04-26	0	0	0	0	0	0	0
5	1844505072	2016-04-27	4	0	0	0	0	0	0
6	1844505072	2016-05-02	0	0	0	0	0	0	0
7	1844505072	2016-05-07	0	0	0	0	0	0	0
8	1844505072	2016-05-08	0	0	0	0	0	0	0
9	1844505072	2016-05-09	0	0	0	0	0	0	0
10	1844505072	2016-05-10	0	0	0	0	0	0	0

```
# ... with 68 more rows, 6 more variables: sedentary_active_distance <dbl>,
# very_active_minutes <dbl>, fairly_active_minutes <dbl>,
# lightly_active_minutes <dbl>, sedentary_minutes <dbl>, calories <dbl>, and
# abbreviated variable names 1: activity_date, 2: total_steps,
# 3: total_distance, 4: tracker_distance, 5: logged_activities_distance,
# 6: very_active_distance, 7: moderately_active_distance,
# 8: light_active_distance
```

From our inspection above, we can see that we just need to delete the entries where total_steps is zero and will take care of the rest.

```
daily_activity_clean <-
  filter(daily_activity,
         total_steps != 0,
         total_distance != 0,
         calories != 0)
daily_activity_clean
```

```
# A tibble: 862 x 15
```

	id	activity~1	total~2	total~3	track~4	logge~5	very_~6	moder~7	light~8
	<chr>	<date>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	1503960366	2016-04-12	13162	8.5	8.5	0	1.88	0.550	6.06
2	1503960366	2016-04-13	10735	6.97	6.97	0	1.57	0.690	4.71
3	1503960366	2016-04-14	10460	6.74	6.74	0	2.44	0.400	3.91
4	1503960366	2016-04-15	9762	6.28	6.28	0	2.14	1.26	2.83
5	1503960366	2016-04-16	12669	8.16	8.16	0	2.71	0.410	5.04
6	1503960366	2016-04-17	9705	6.48	6.48	0	3.19	0.780	2.51
7	1503960366	2016-04-18	13019	8.59	8.59	0	3.25	0.640	4.71
8	1503960366	2016-04-19	15506	9.88	9.88	0	3.53	1.32	5.03
9	1503960366	2016-04-20	10544	6.68	6.68	0	1.96	0.480	4.24
10	1503960366	2016-04-21	9819	6.34	6.34	0	1.34	0.350	4.65

```
# ... with 852 more rows, 6 more variables: sedentary_active_distance <dbl>,
# very_active_minutes <dbl>, fairly_active_minutes <dbl>,
# lightly_active_minutes <dbl>, sedentary_minutes <dbl>, calories <dbl>, and
# abbreviated variable names 1: activity_date, 2: total_steps,
# 3: total_distance, 4: tracker_distance, 5: logged_activities_distance,
# 6: very_active_distance, 7: moderately_active_distance,
# 8: light_active_distance
```

```
names(daily_activity)
```

```
[1] "id" "activity_date"
[3] "total_steps" "total_distance"
[5] "tracker_distance" "logged_activities_distance"
[7] "very_active_distance" "moderately_active_distance"
[9] "light_active_distance" "sedentary_active_distance"
[11] "very_active_minutes" "fairly_active_minutes"
[13] "lightly_active_minutes" "sedentary_minutes"
[15] "calories"
```

```
# Check the attributes again
```

```
cat("Before deleting the entries\n\n")
```

Before deleting the entries

```
select(daily_activity, total_steps, total_distance, calories) %>%
  summary()
```

total_steps	total_distance	calories
Min. : 0	Min. : 0.000	Min. : 0
1st Qu.: 3790	1st Qu.: 2.620	1st Qu.: 1828
Median : 7406	Median : 5.245	Median : 2134
Mean : 7638	Mean : 5.490	Mean : 2304
3rd Qu.: 10727	3rd Qu.: 7.713	3rd Qu.: 2793
Max. : 36019	Max. : 28.030	Max. : 4900

```
cat("\n\n\n",
    "\t\t vs",
    "\n\n\n")
```

vs

```
cat("After deleting the entries\n\n")
```

After deleting the entries

```
select(daily_activity_clean, total_steps, total_distance, calories) %>%
  summary()
```

total_steps	total_distance	calories
Min. : 8	Min. : 0.010	Min. : 52
1st Qu.: 4927	1st Qu.: 3.373	1st Qu.: 1857
Median : 8054	Median : 5.590	Median : 2220
Mean : 8329	Mean : 5.986	Mean : 2362
3rd Qu.: 11096	3rd Qu.: 7.905	3rd Qu.: 2832
Max. : 36019	Max. : 28.030	Max. : 4900

We can see that the observation we removed affected our mean and median.

Clean the daily_sleep dataset

```
# Check daily_sleep dataset before cleaning
glimpse(daily_sleep)
```

```

Rows: 413
Columns: 5
$ Id          <dbl> 1503960366, 1503960366, 1503960366, 1503960366, 150~
$ SleepDay    <chr> "4/12/2016 12:00:00 AM", "4/13/2016 12:00:00 AM", "~
$ TotalSleepRecords <dbl> 1, 2, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ~
$ TotalMinutesAsleep <dbl> 327, 384, 412, 340, 700, 304, 360, 325, 361, 430, 2~
$ TotalTimeInBed  <dbl> 346, 407, 442, 367, 712, 320, 377, 364, 384, 449, 3~

```

```
# Check missing values and duplicates
```

```

cat("\n",
    "Missing values:",
    sum(is.na(daily_sleep)),
    "\n",
    "Duplicate values:",
    sum(duplicated(daily_sleep)),
    "\n",
    "Unique Ids:",
    n_distinct(daily_sleep$Id)
)

```

Missing values: 0

Duplicate values: 3

Unique Ids: 24

Let us clean:

- Change column names to lower case because R is case sensitive
- Change “Id” from double to a character because the number represents a category
- Change “SleepDay” from char to date. Since the time component of this column is the same for each observation “12:00:00 AM”, we can remove it. This will help us merge this dataset with daily_activity later
- Delete duplicates (3 observations are duplicates)

```
# Clean daily_sleep dataset
```

```

daily_sleep_clean <-
  # Clean column names
  clean_names(daily_sleep) %>%
  # Correct column types
  mutate(id = as.character(id)) %>% # from double to chr
  mutate(sleep_day = as.Date(sleep_day,
                             format = "%m/%d/%Y")) %>% # from chr to date

  # Remove duplicate rows
  distinct()

```

```
# Check clean daily_sleep dataset
```

```
glimpse(daily_sleep_clean)
```

```

Rows: 410
Columns: 5
$ id          <chr> "1503960366", "1503960366", "1503960366", "150396~
$ sleep_day    <date> 2016-04-12, 2016-04-13, 2016-04-15, 2016-04-16, ~
$ total_sleep_records <dbl> 1, 2, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
$ total_minutes_asleep <dbl> 327, 384, 412, 340, 700, 304, 360, 325, 361, 430, ~

```



```
$ total_time_in_bed      <dbl> 346, 407, 442, 367, 712, 320, 377, 364, 384, 449,~
```

```
# Check missing values and duplicates after cleaning
```

```
cat("\n",  
    "Missing values:",  
    sum(is.na(daily_sleep_clean)),  
    "\n",  
    "Duplicate values:",  
    sum(duplicated(daily_sleep_clean)))
```

```
Missing values: 0
```

```
Duplicate values: 0
```

Clean the hourly datasets (hourly_calories, hourly_intensities, and hourly_steps)

```
# Check hourly_calories dataset before cleaning
```

```
glimpse(hourly_calories)
```

```
Rows: 22,099
```

```
Columns: 3
```

```
$ Id      <dbl> 1503960366, 1503960366, 1503960366, 1503960366, 150396036~
```

```
$ ActivityHour <chr> "4/12/2016 12:00:00 AM", "4/12/2016 1:00:00 AM", "4/12/20~
```

```
$ Calories    <dbl> 81, 61, 59, 47, 48, 48, 48, 47, 68, 141, 99, 76, 73, 66, ~
```

```
# Check missing values and duplicates
```

```
cat("\n",  
    "Missing values:",  
    sum(is.na(hourly_calories)),  
    "\n",  
    "Duplicate values:",  
    sum(duplicated(hourly_calories)))
```

```
Missing values: 0
```

```
Duplicate values: 0
```

```
# Check hourly_intensities dataset before cleaning
```

```
glimpse(hourly_intensities)
```

```
Rows: 22,099
```

```
Columns: 4
```

```
$ Id      <dbl> 1503960366, 1503960366, 1503960366, 1503960366, 15039~
```

```
$ ActivityHour <chr> "4/12/2016 12:00:00 AM", "4/12/2016 1:00:00 AM", "4/1~
```

```
$ TotalIntensity <dbl> 20, 8, 7, 0, 0, 0, 0, 0, 13, 30, 29, 12, 11, 6, 36, 5~
```

```
$ AverageIntensity <dbl> 0.333333, 0.133333, 0.116667, 0.000000, 0.000000, 0.0~
```

```
# Check missing values and duplicates
```

```
cat("\n",  
    "Missing values:",  
    sum(is.na(hourly_intensities)),  
    "\n",  
    "Duplicate values:",  
    sum(duplicated(hourly_intensities)))
```

```
Missing values: 0
```

Duplicate values: 0

```
# Check hourly_steps dataset before cleaning
glimpse(hourly_steps)
```

Rows: 22,099

Columns: 3

```
$ Id          <dbl> 1503960366, 1503960366, 1503960366, 1503960366, 150396036~
$ ActivityHour <chr> "4/12/2016 12:00:00 AM", "4/12/2016 1:00:00 AM", "4/12/20~
$ StepTotal    <dbl> 373, 160, 151, 0, 0, 0, 0, 0, 250, 1864, 676, 360, 253, 2~
```

```
# Check missing values and duplicates
```

```
cat("\n",
    "Missing values:",
    sum(is.na(hourly_steps)),
    "\n",
    "Duplicate values:",
    sum(duplicated(hourly_steps)))
```

Missing values: 0

Duplicate values: 0

Join hourly datasets to create a hourly_actitvity dataset

These datasets shared the same Id and Activity_hour, let us join them into a new dataset (hourly_activity) before we clean them.

```
# Join the hourly datasets (hourly_calories, hourly_intensities, and hourly_steps)
```

```
hourly_activity <-
  inner_join(hourly_calories,
             hourly_intensities,
             by = c("Id", "ActivityHour"))
```

```
hourly_activity <-
  inner_join(hourly_activity, hourly_steps, by = c("Id", "ActivityHour"))
```

```
# Check hourly_activity dataset before cleaning
```

```
glimpse(hourly_activity)
```

Rows: 22,099

Columns: 6

```
$ Id          <dbl> 1503960366, 1503960366, 1503960366, 1503960366, 15039~
$ ActivityHour <chr> "4/12/2016 12:00:00 AM", "4/12/2016 1:00:00 AM", "4/1~
$ Calories     <dbl> 81, 61, 59, 47, 48, 48, 48, 47, 68, 141, 99, 76, 73, ~
$ TotalIntensity <dbl> 20, 8, 7, 0, 0, 0, 0, 0, 13, 30, 29, 12, 11, 6, 36, 5~
$ AverageIntensity <dbl> 0.333333, 0.133333, 0.116667, 0.000000, 0.000000, 0.0~
$ StepTotal    <dbl> 373, 160, 151, 0, 0, 0, 0, 0, 250, 1864, 676, 360, 25~
```

```
# Check missing values and duplicates
```

```
cat("\n",
    "Missing values:",
    sum(is.na(hourly_activity)),
    "\n",
    "Duplicate values:",
    sum(duplicated(hourly_activity)))
```

Missing values: 0
Duplicate values: 0

Let us clean:

- Change column names to lower case because R is case sensitive
- Change “Id” from double to a character because the number represents a category
- Change “ActivityHour” from char to datetime

Note: The default timezone is UTC.

```
# Clean hourly_activity dataset

hourly_activity_clean <-
  # Clean column names
  clean_names(hourly_activity) %>%
  # Correct column types
  mutate(id = as.character(id)) %>% # from double to chr
  mutate(activity_hour = as_datetime(activity_hour,
                                     format = "%m/%d/%Y %I:%M:%S %p")) %>% # from chr to datetime

  # Remove duplicate rows
  distinct()

# Check clean daily_activity dataset
glimpse(hourly_activity_clean)
```

Rows: 22,099

Columns: 6

```
$ id          <chr> "1503960366", "1503960366", "1503960366", "150396036~
$ activity_hour <dtm> 2016-04-12 00:00:00, 2016-04-12 01:00:00, 2016-04-1~
$ calories     <dbl> 81, 61, 59, 47, 48, 48, 48, 47, 68, 141, 99, 76, 73, ~
$ total_intensity <dbl> 20, 8, 7, 0, 0, 0, 0, 0, 13, 30, 29, 12, 11, 6, 36, ~
$ average_intensity <dbl> 0.333333, 0.133333, 0.116667, 0.000000, 0.000000, 0.~
$ step_total    <dbl> 373, 160, 151, 0, 0, 0, 0, 0, 250, 1864, 676, 360, 2~
```

```
# Check missing values and duplicates after cleaning
cat("\n",
    "Missing values:",
    sum(is.na(hourly_activity_clean)),
    "\n",
    "Duplicate values:",
    sum(duplicated(hourly_activity_clean)))
```

Missing values: 0
Duplicate values: 0

```
# as_datetime() converts with default timezone = "UTC"
```

Clean the minute_sleep dataset

```
# Check minute_sleep dataset before cleaning
glimpse(minute_sleep)
```

Rows: 188,521

Columns: 4

```
$ Id      <dbl> 1503960366, 1503960366, 1503960366, 1503960366, 1503960366, 1503~
$ date    <chr> "4/12/2016 2:47:30 AM", "4/12/2016 2:48:30 AM", "4/12/2016 2:49:~
$ value    <dbl> 3, 2, 1, 1, 1, 1, 1, 2, 2, 2, 3, 3, 3, 3, 3, 2, 1, 1, 1, 1, 1, 1~
$ logId    <dbl> 11380564589, 11380564589, 11380564589, 11380564589, 11380564589,~
```

```
# Check missing values and duplicates
```

```
cat("\n",
    "Missing values:",
    sum(is.na(minute_sleep)),
    "\n",
    "Duplicate values:",
    sum(duplicated(minute_sleep)),
    "\n",
    "Unique Ids:",
    n_distinct(minute_sleep$Id))
```

Missing values: 0

Duplicate values: 543

Unique Ids: 24

Let us clean:

- Change column names to lower case because R is case sensitive.
- Change “Id” from double to a character because the number represents a category.
- Change “date” from char to datetime.
- Change “value” from double to factor. Value indicates the sleep state: 1 = asleep, 2 = restless, 3 = awake. For details see: [Fitbit data dictionary](#)
- Remove duplicate values: 543.

```
# Clean minute_sleep dataset
```

```
minute_sleep_clean <-
  # Clean column names
  clean_names(minute_sleep) %>%
  # Correct column types
  mutate(value = as.factor(value)) %>% # from double to chr
  mutate(id = as.character(id)) %>% # from double to chr
  mutate(date = as_datetime(date,
                             format = "%m/%d/%Y %I:%M:%S %p")) %>% # From chr to datetime

  # Remove duplicate rows
  distinct()
```

```
# Check clean daily_activity dataset
```

```
glimpse(minute_sleep_clean)
```

Rows: 187,978

Columns: 4

```
$ id      <chr> "1503960366", "1503960366", "1503960366", "1503960366", "150396~
$ date     <dtm> 2016-04-12 02:47:30, 2016-04-12 02:48:30, 2016-04-12 02:49:30,~
$ value     <fct> 3, 2, 1, 1, 1, 1, 1, 2, 2, 2, 3, 3, 3, 3, 3, 2, 1, 1, 1, 1, 1, ~
$ log_id    <dbl> 11380564589, 11380564589, 11380564589, 11380564589, 11380564589~
```

```
# Check missing values and duplicates after cleaning
```

```
cat("\n",
```

```

"Missing values:",
sum(is.na(minute_sleep_clean)),
"\n",
"Duplicate values:",
sum(duplicated(minute_sleep_clean)))

```

Missing values: 0
Duplicate values: 0

Clean the seconds_heartrate dataset

```

# Check seconds_heartrate set before cleaning
glimpse(seconds_heartrate)

```

```

Rows: 2,483,658
Columns: 3
$ Id      <dbl> 2022484408, 2022484408, 2022484408, 2022484408, 2022484408, 2022~
$ Time    <chr> "4/12/2016 7:21:00 AM", "4/12/2016 7:21:05 AM", "4/12/2016 7:21:~
$ Value   <dbl> 97, 102, 105, 103, 101, 95, 91, 93, 94, 93, 92, 89, 83, 61, 60, ~

```

```

# Check missing values and duplicates
cat(
  "\n",
  "Missing values:", sum(is.na(seconds_heartrate)),
  "\n",
  "Duplicate values:", sum(duplicated(seconds_heartrate))
)

```

Missing values: 0
Duplicate values: 0

Let us clean:

- Change column names to lower case because R is case sensitive
- Change “Id” from double to a character because the number represents a category
- Change “Time” from char to datetime and rename it date_time
- Rename “Value” to heart_rate [Fitbit data dictionary](#)

```

# Clean seconds_heartrate dataset

seconds_heartrate_clean <-
  # Clean column names
  clean_names(seconds_heartrate) %>%
  # Correct column types
  mutate(id = as.character(id)) %>% # from double to chr
  mutate(time = as_datetime(time,
                             format = "%m/%d/%Y %I:%M:%S %p")) %>% # from chr to datetime

  # Rename columns
  rename(date_time = time,
         heart_rate = value) %>%
  # Remove duplicate rows
  distinct()

# Check clean daily_activity dataset

```

```
glimpse(seconds_hearttrate_clean)
```

Rows: 2,483,658

Columns: 3

```
$ id      <chr> "2022484408", "2022484408", "2022484408", "2022484408", "20~
$ date_time <dtm> 2016-04-12 07:21:00, 2016-04-12 07:21:05, 2016-04-12 07:21~
$ heart_rate <dbl> 97, 102, 105, 103, 101, 95, 91, 93, 94, 93, 92, 89, 83, 61,~
```

```
# Check missing values and duplicates after cleaning
```

```
cat("\n",
    "Missing values:",
    sum(is.na(seconds_hearttrate_clean)),
    "\n",
    "Duplicate values:",
    sum(duplicated(seconds_hearttrate_clean)))
```

Missing values: 0

Duplicate values: 0

```
# as_datetime() converts with default timezone = "UTC"
```

Clean the weight_logs dataset

```
# Check weight_logs set before cleaning
```

```
glimpse(weight_logs)
```

Rows: 67

Columns: 8

```
$ Id      <dbl> 1503960366, 1503960366, 1927972279, 2873212765, 2873212~
$ Date    <chr> "5/2/2016 11:59:59 PM", "5/3/2016 11:59:59 PM", "4/13/2~
$ WeightKg <dbl> 52.6, 52.6, 133.5, 56.7, 57.3, 72.4, 72.3, 69.7, 70.3, ~
$ WeightPounds <dbl> 115.9631, 115.9631, 294.3171, 125.0021, 126.3249, 159.6~
$ Fat      <dbl> 22, NA, NA, NA, NA, 25, NA, NA, NA, NA, NA, NA, NA, NA,~
$ BMI      <dbl> 22.65, 22.65, 47.54, 21.45, 21.69, 27.45, 27.38, 27.25,~
$ IsManualReport <lgl> TRUE, TRUE, FALSE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, ~
$ LogId    <dbl> 1.462234e+12, 1.462320e+12, 1.460510e+12, 1.461283e+12,~
```

```
# Check missing values and duplicates
```

```
cat("\n",
    "Missing values:",
    sum(is.na(weight_logs)),
    "\n",
    "Duplicate values:",
    sum(duplicated(weight_logs)))
```

Missing values: 65

Duplicate values: 0

Let us clean: - Change column names to lower case because R is case sensitive.

- Change “Id” from double to a character because the number represents a category
- Change “Date” from char to datetime and rename it date_time.

- Change NA to 0 in the column “fat.”

```
# Clean weight_logs dataset

weight_logs_clean <-
  # Clean column names
  clean_names(weight_logs) %>%
  # Correct column types
  mutate(id = as.character(id)) %>% # from double to chr
  mutate(date = as_datetime(date,
                             format = "%m/%d/%Y %I:%M:%S %p")) %>% # from chr to datetime

  # Rename columns
  rename(date_time = date) %>%
  # Remove duplicate rows
  distinct()

# Change NA to 0 in the column "fat"
weight_logs_clean$fat[is.na(weight_logs_clean$fat)] <- 0

# Check clean daily_activity dataset
glimpse(weight_logs_clean)
```

```
Rows: 67
Columns: 8
$ id          <chr> "1503960366", "1503960366", "1927972279", "2873212765~
$ date_time   <dtm> 2016-05-02 23:59:59, 2016-05-03 23:59:59, 2016-04-13~
$ weight_kg    <dbl> 52.6, 52.6, 133.5, 56.7, 57.3, 72.4, 72.3, 69.7, 70.3~
$ weight_pounds <dbl> 115.9631, 115.9631, 294.3171, 125.0021, 126.3249, 159~
$ fat          <dbl> 22, 0, 0, 0, 0, 25, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ~
$ bmi          <dbl> 22.65, 22.65, 47.54, 21.45, 21.69, 27.45, 27.38, 27.2~
$ is_manual_report <lgl> TRUE, TRUE, FALSE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE~
$ log_id       <dbl> 1.462234e+12, 1.462320e+12, 1.460510e+12, 1.461283e+1~
```

```
# Check missing values and duplicates after cleaning
```

```
cat("\n",
    "Missing values:",
    sum(is.na(weight_logs_clean)),
    "\n",
    "Duplicate values:",
    sum(duplicated(weight_logs_clean)))
```

Missing values: 0

Duplicate values: 0

Distribution of ids across datasets

```
# Loop through each dataset and print th number o funiqu ids
datasets <- c(
  "daily_activity_clean",
  "daily_sleep_clean",
  "hourly_activity_clean",
```

```

    "minute_sleep_clean",
    "seconds_heartrate_clean",
    "weight_logs_clean"
)

results_df <- data.frame(Dataset = character(0), distinct_IDs = integer(0))

for (dataset_name in datasets) {
  dataset <- get(dataset_name) # Retrieve the dataset by its name
  distinct_ids <- length(unique(dataset$id)) # Calculate the number of distinct IDs

  result_row <- data.frame(Dataset = dataset_name, distinct_IDs = distinct_ids)
  results_df <- bind_rows(results_df, result_row)
}

sorted_results <- results_df %>% arrange(- distinct_IDs )

print(sorted_results)

```

	Dataset	distinct_IDs
1	daily_activity_clean	33
2	hourly_activity_clean	33
3	daily_sleep_clean	24
4	minute_sleep_clean	24
5	seconds_heartrate_clean	14
6	weight_logs_clean	8

- Differences in the number of unique IDs between the datasets can imply discrepancies in data collection methods, data incompleteness, or differing levels of user engagement.

Export clean datasets

```

# To uncomment the following code, select all the lines and press shift + control + c on Mac

# write.csv(daily_activity_clean,
#           "daily_activity_clean.csv",
#           row.names = FALSE)
#
# write.csv(daily_sleep_clean,
#           "daily_sleep_clean.csv",
#           row.names = FALSE)
#
# write.csv(hourly_activity_clean,
#           "hourly_activity_clean.csv",
#           row.names = FALSE)
#
# write.csv(minute_sleep_clean,
#           "minute_sleep_clean.csv",
#           row.names = FALSE)
#
# write.csv(seconds_heartrate_clean,
#           "seconds_heartrate_clean.csv",

```



```
#         row.names = FALSE)
#
# write.csv(weight_logs_clean ,
#           "weight_logs_clean .csv",
#           row.names = FALSE)
```

Analyze phase: Exploratory data analysis

EDA for daily_activity_clean

```
str(daily_activity_clean)

tibble [862 x 15] (S3: tbl_df/tbl/data.frame)
 $ id                : chr [1:862] "1503960366" "1503960366" "1503960366" "1503960366" ...
 $ activity_date      : Date[1:862], format: "2016-04-12" "2016-04-13" ...
 $ total_steps        : num [1:862] 13162 10735 10460 9762 12669 ...
 $ total_distance     : num [1:862] 8.5 6.97 6.74 6.28 8.16 ...
 $ tracker_distance   : num [1:862] 8.5 6.97 6.74 6.28 8.16 ...
 $ logged_activities_distance: num [1:862] 0 0 0 0 0 0 0 0 0 0 ...
 $ very_active_distance : num [1:862] 1.88 1.57 2.44 2.14 2.71 ...
 $ moderately_active_distance: num [1:862] 0.55 0.69 0.4 1.26 0.41 ...
 $ light_active_distance : num [1:862] 6.06 4.71 3.91 2.83 5.04 ...
 $ sedentary_active_distance : num [1:862] 0 0 0 0 0 0 0 0 0 0 ...
 $ very_active_minutes  : num [1:862] 25 21 30 29 36 38 42 50 28 19 ...
 $ fairly_active_minutes : num [1:862] 13 19 11 34 10 20 16 31 12 8 ...
 $ lightly_active_minutes : num [1:862] 328 217 181 209 221 164 233 264 205 211 ...
 $ sedentary_minutes    : num [1:862] 728 776 1218 726 773 ...
 $ calories            : num [1:862] 1985 1797 1776 1745 1863 ...
```

Univariate analysis for daily_activity_clean

```
# Subset numeric columns
num_df <- select_if(daily_activity_clean, is.numeric)

# Identify numeric columns
colnames(num_df)
```

Numerical variables

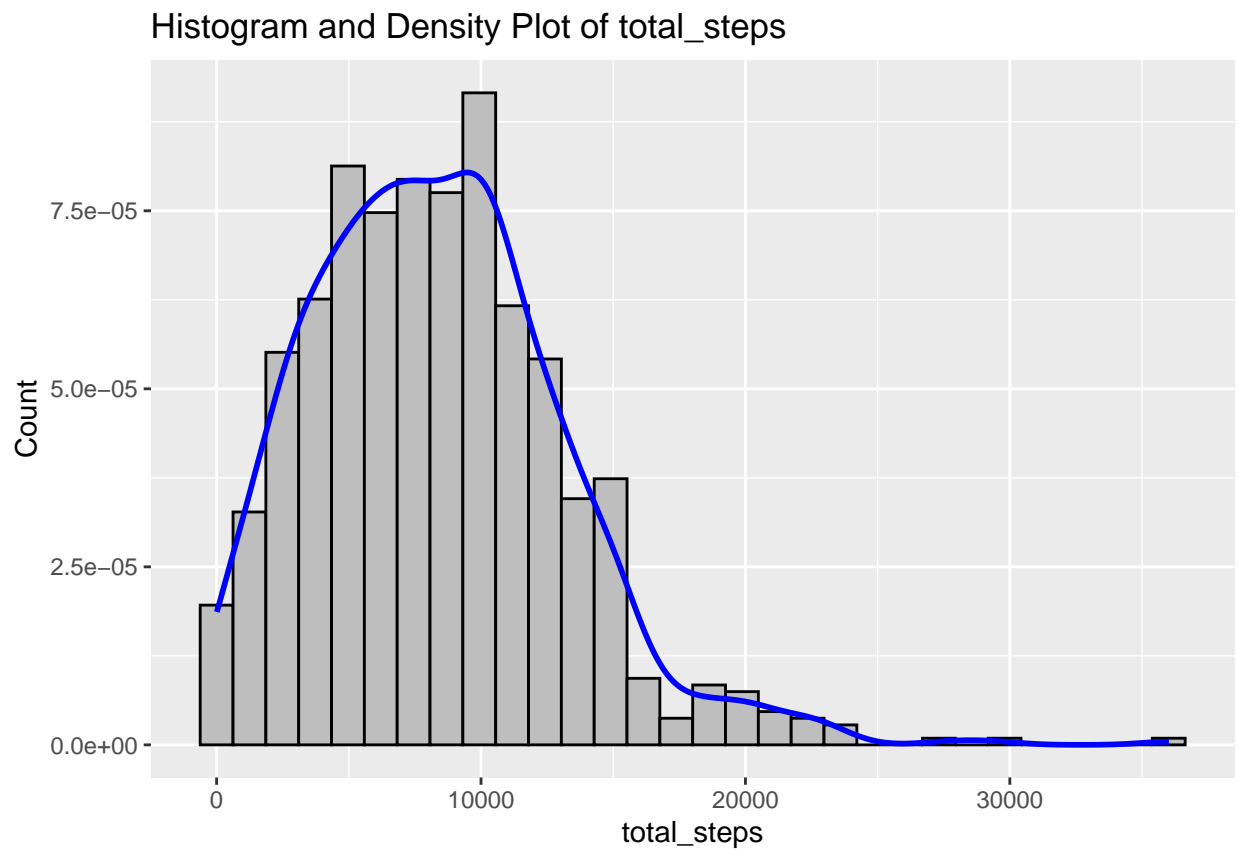
```
[1] "total_steps"           "total_distance"
[3] "tracker_distance"     "logged_activities_distance"
[5] "very_active_distance"  "moderately_active_distance"
[7] "light_active_distance" "sedentary_active_distance"
[9] "very_active_minutes"   "fairly_active_minutes"
[11] "lightly_active_minutes" "sedentary_minutes"
[13] "calories"

# plotting all numerical variables
col_names <- colnames(num_df)
for (i in col_names) {
  suppressWarnings(print(
    ggplot(num_df, aes(num_df[[i]])) +
    geom_histogram(
```

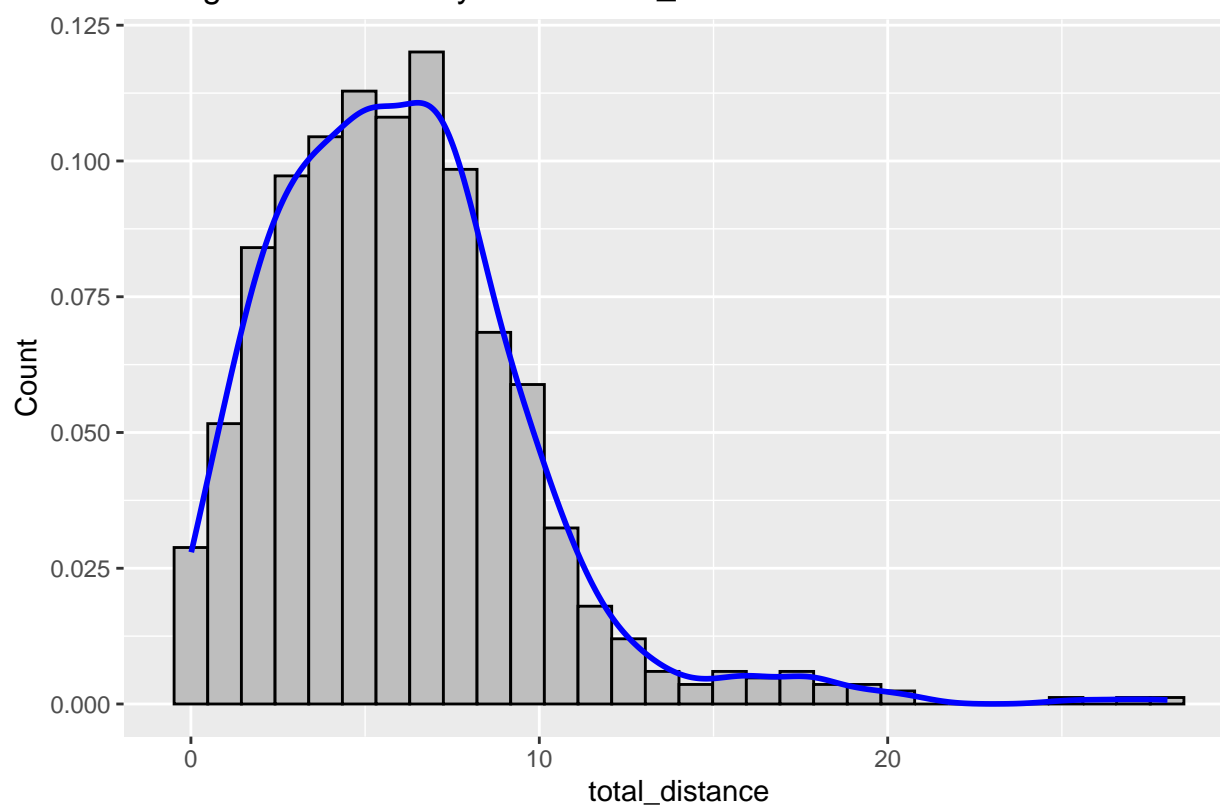
```

    bins = 30,
    color = "black",
    fill = "gray",
    aes(y = ..density..)
  ) +
  geom_density(
    color = "blue",
    size = 1
  ) +
  xlab(i) + ylab("Count") +
  ggtitle(paste("Histogram and Density Plot of", i))
))
}


```



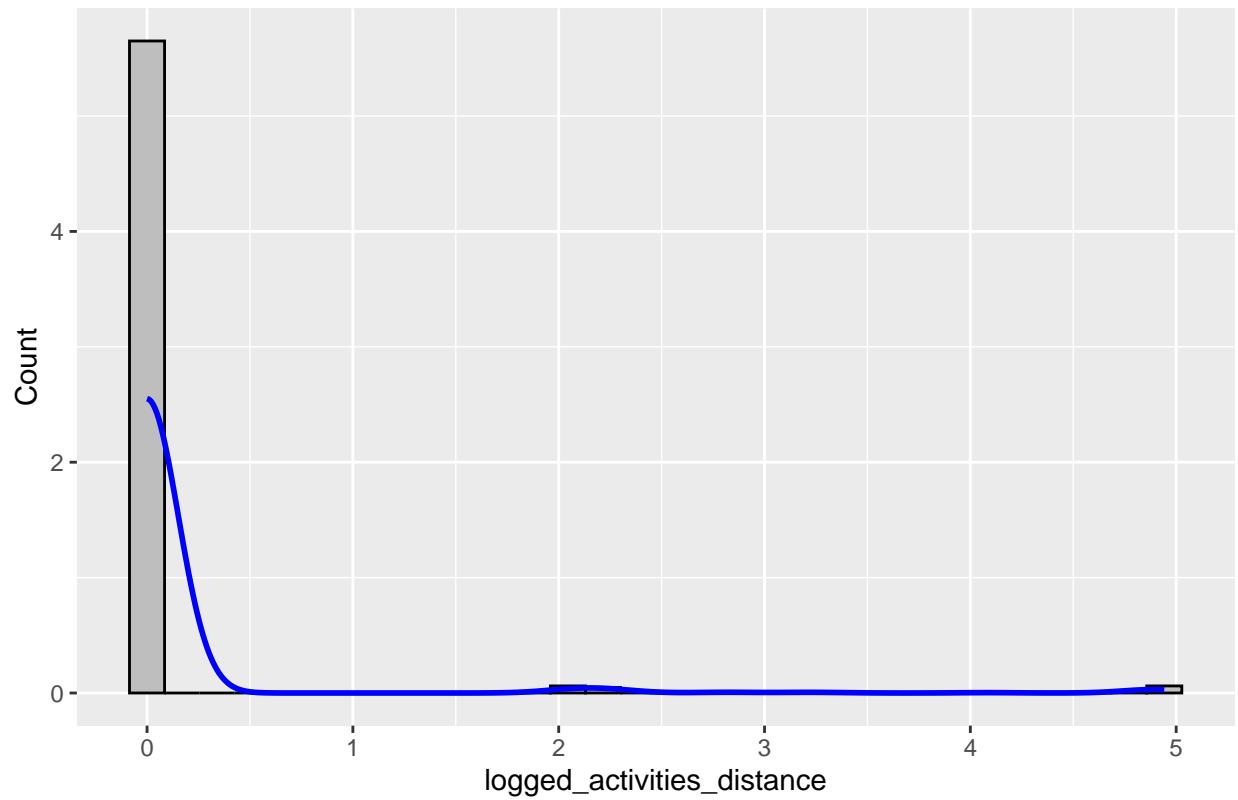
Histogram and Density Plot of total_distance



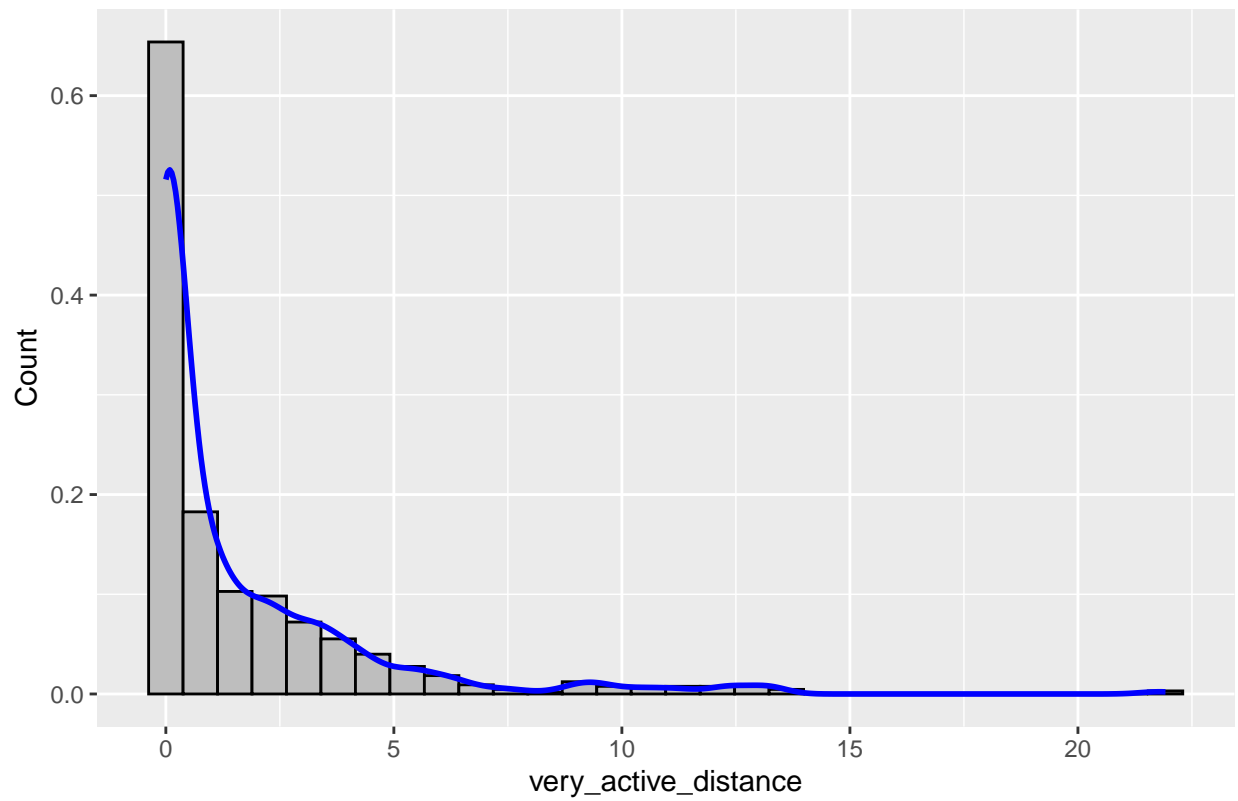
Histogram and Density Plot of tracker_distance



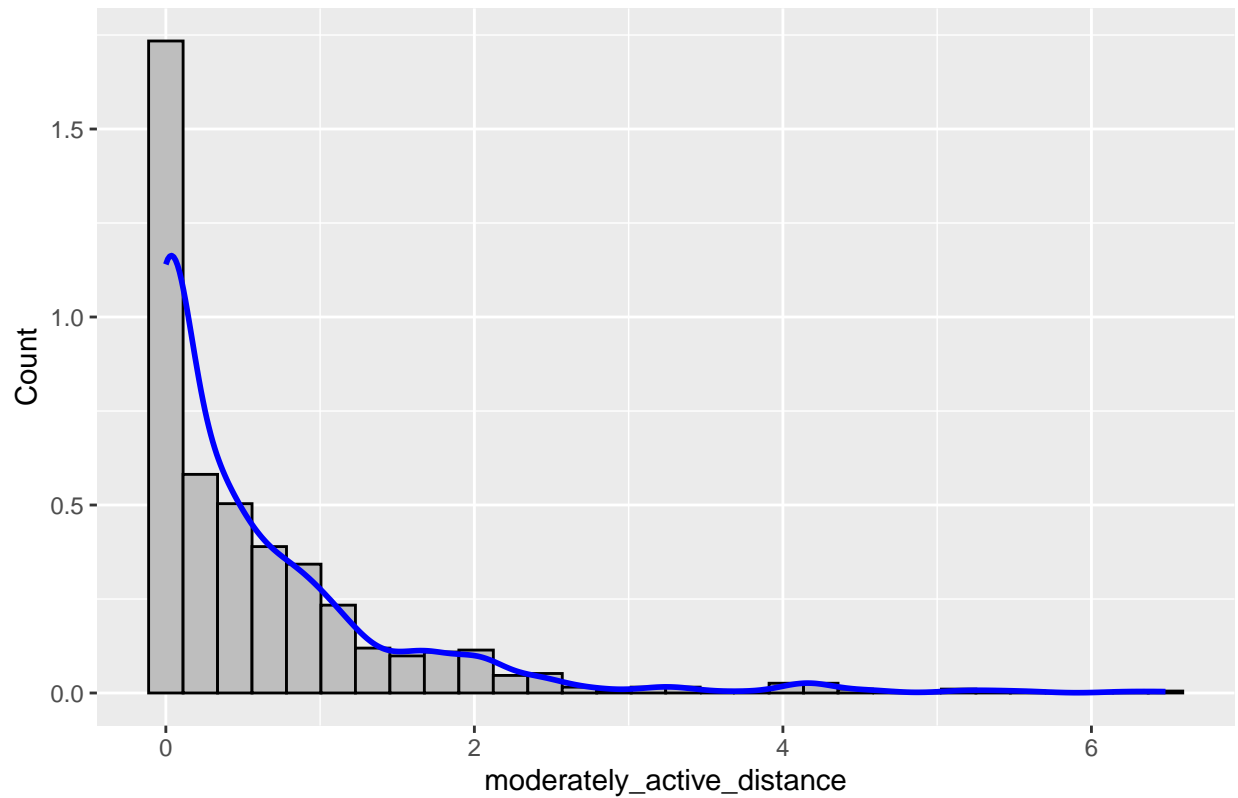
Histogram and Density Plot of logged_activities_distance




Histogram and Density Plot of very_active_distance



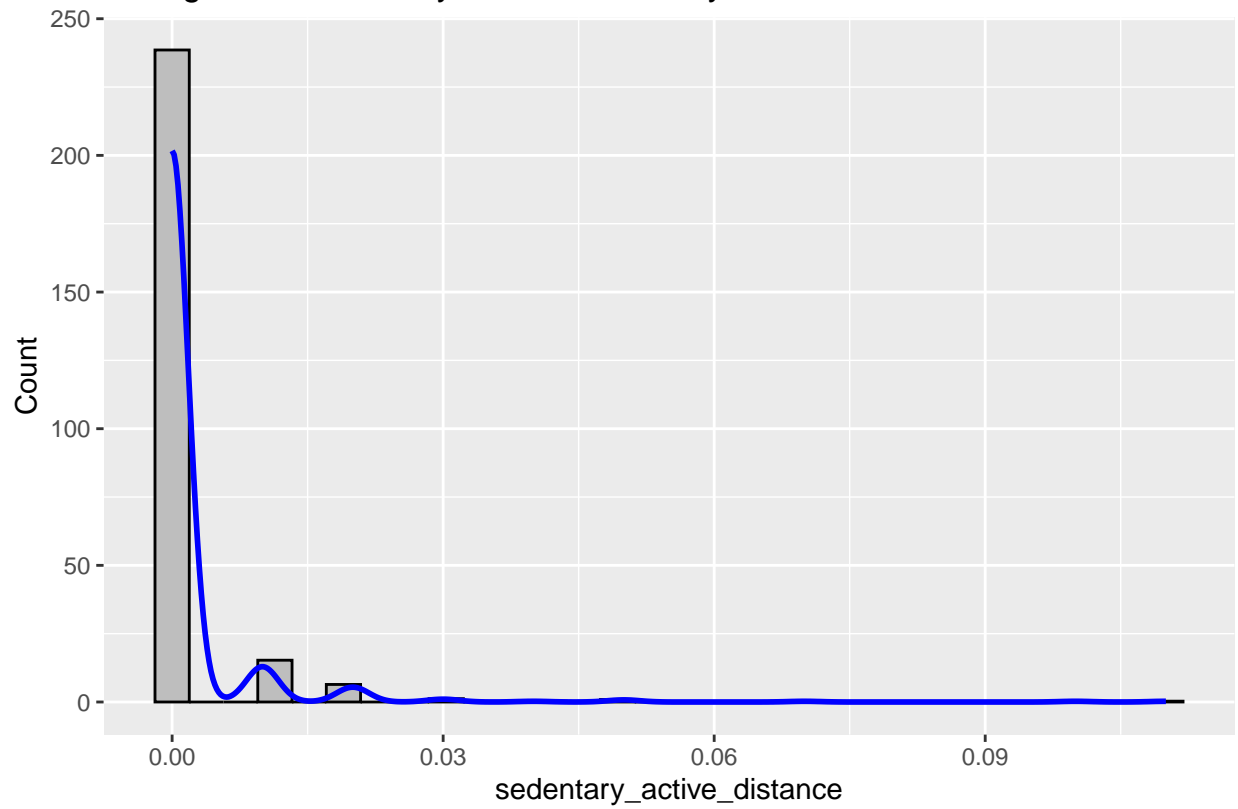
Histogram and Density Plot of moderately_active_distance




Histogram and Density Plot of light_active_distance

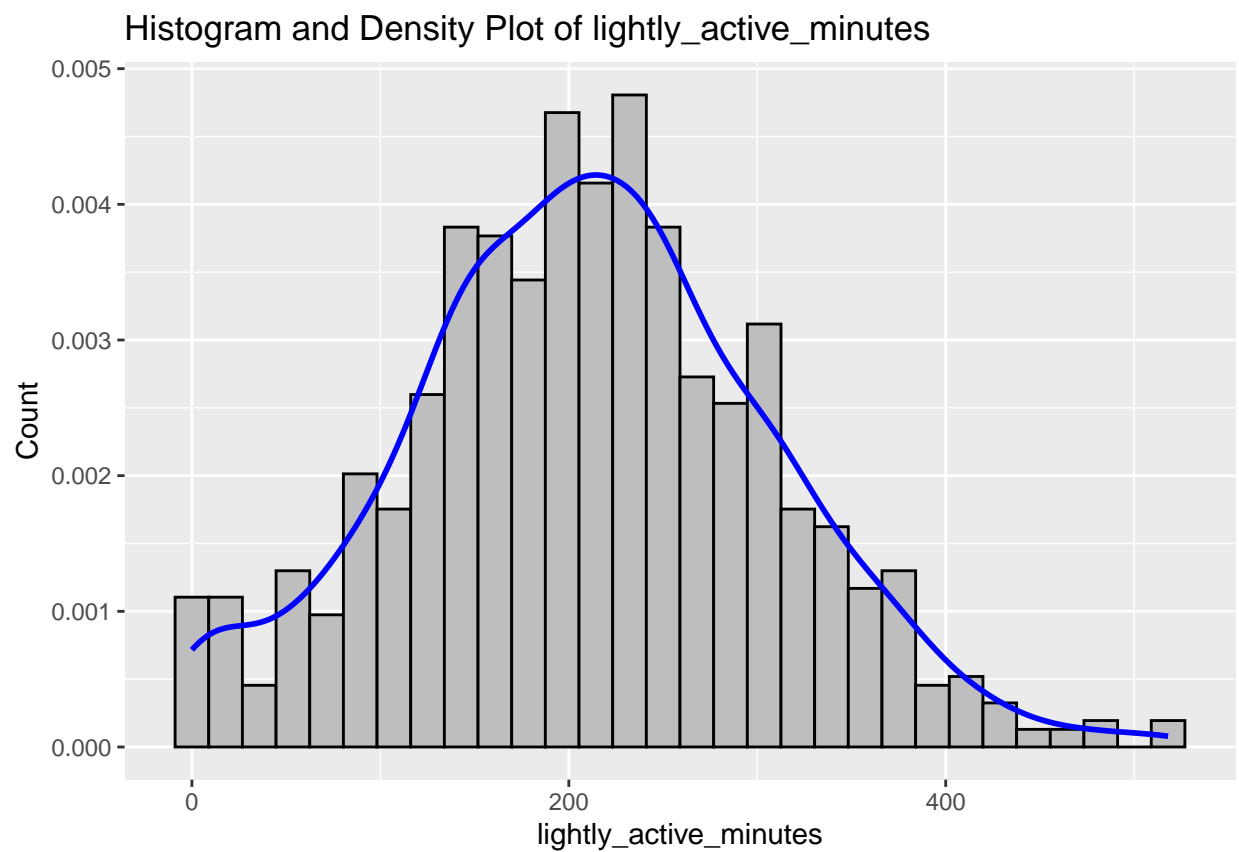
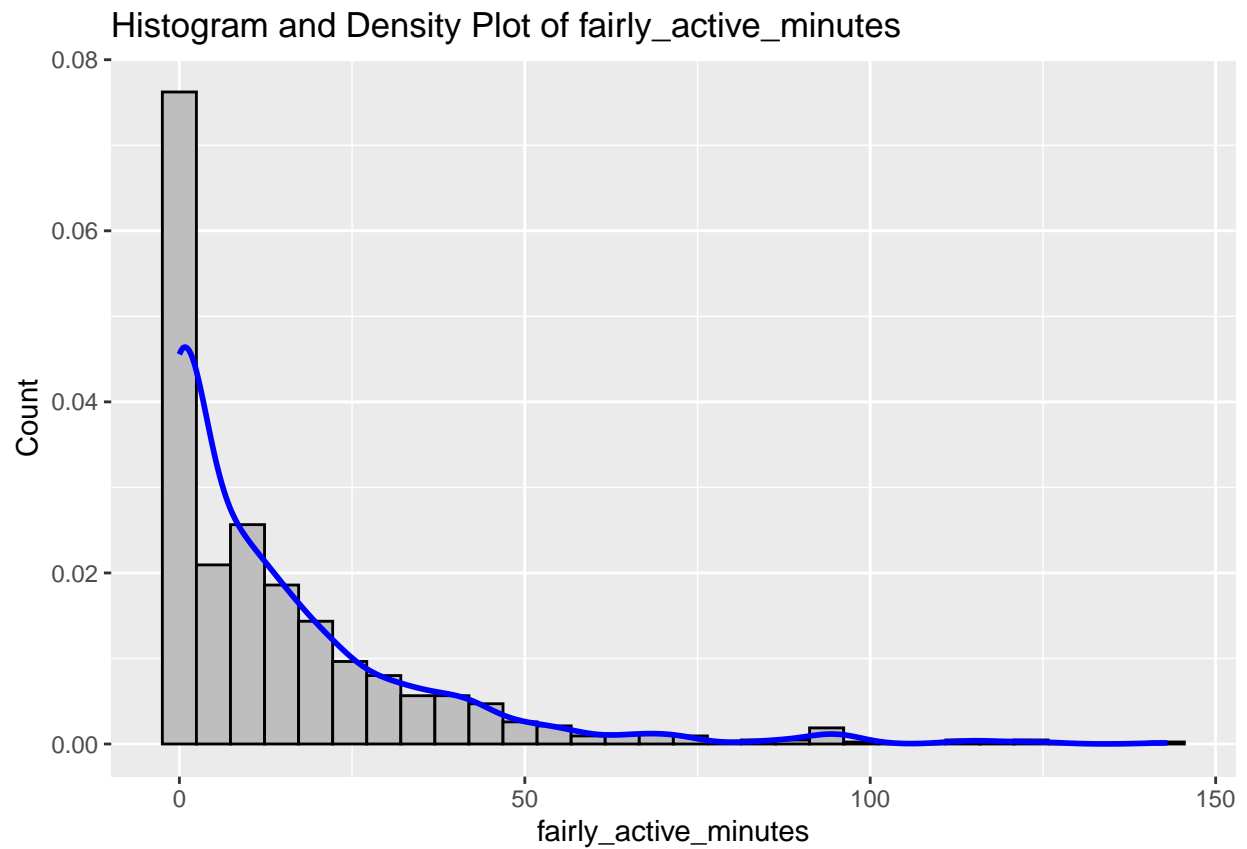


Histogram and Density Plot of sedentary_active_distance

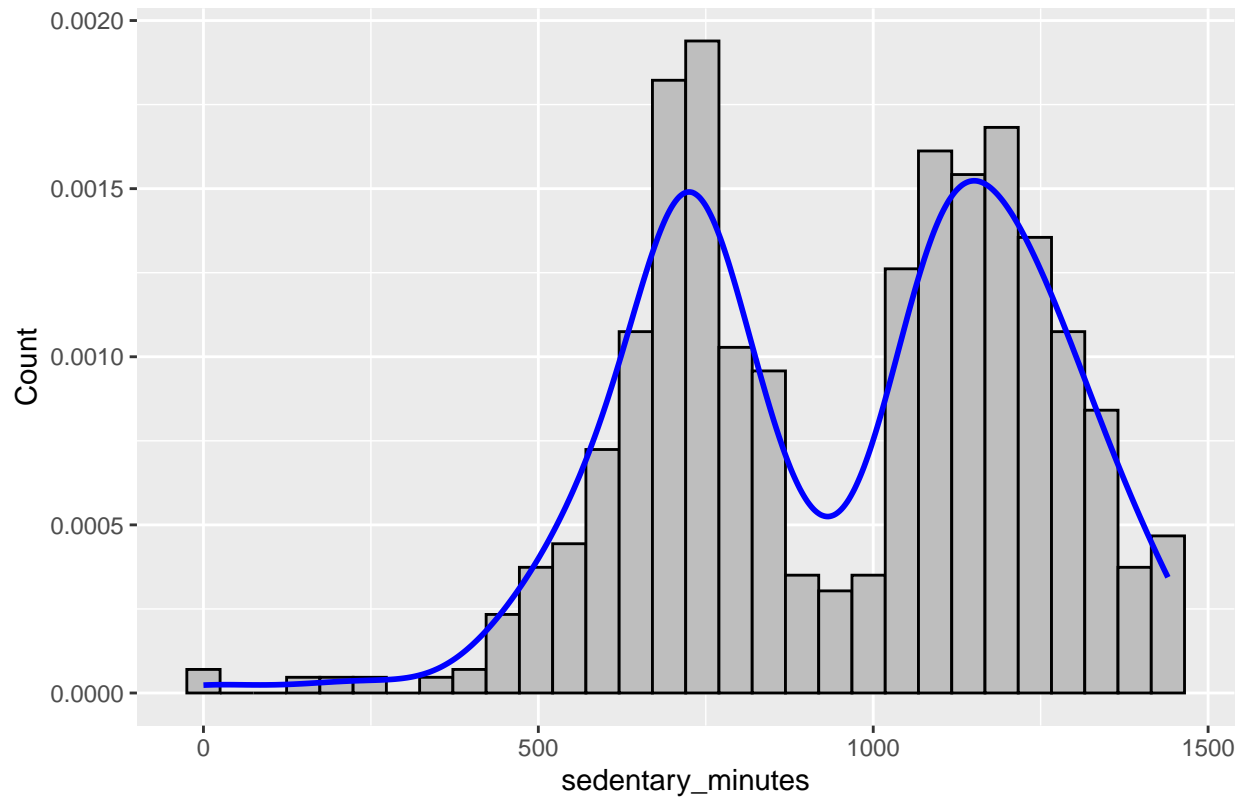


Histogram and Density Plot of very_active_minutes






Histogram and Density Plot of sedentary_minutes



Histogram and Density Plot of calories



Observations:

- Many variables show a right-skewed distribution: a larger number of data values are located on the left side of the curve.
- The variables `total_steps`, `total_distance`, `tracker_distance` have a similar distribution. We can explore their correlations later.
- Since the distributions are not normal. The median is a better indicator of central tendency for the numerical variables in these dataset.
- **The variable “logged_activities_distance” and “sedentary_active_distance” might not provide useful information since most of the data points are zero. It seems that the users are not logging the distance frequently.**
- The following variables seem related. We will explore them further in the bivariate analysis section:
 - `sedentary_minutes`; `sedentary_active_distance`
 - `lightly_active_minutes`; `light_active_distance`
 - `fairly_active_minutes`; `moderately_active_distance`
 - `very_active_minutes`; `very_active_distance`
- The variables `calories` and `sedentary_minutes` exhibit a multimodal distribution, indicating the presence of subpopulations within the data. In this dataset, gender could be a potential variable that would result in a bimodal distribution when examining histograms of calories and sedentary minutes. Unfortunately, the gender of the users is not provided, limiting our ability to confirm this hypothesis.


```
# Subset numeric columns

select_if(daily_activity_clean, negate(is.numeric))
```

Categorical variables

```
# A tibble: 862 x 2
  id          activity_date
<chr>      <date>
1 1503960366 2016-04-12
2 1503960366 2016-04-13
3 1503960366 2016-04-14
4 1503960366 2016-04-15
5 1503960366 2016-04-16
6 1503960366 2016-04-17
7 1503960366 2016-04-18
8 1503960366 2016-04-19
9 1503960366 2016-04-20
10 1503960366 2016-04-21
# ... with 852 more rows

# Check counts by id
ggplot(data=daily_activity_clean) +
  geom_bar(mapping = aes (x= reorder(id, id,length)))+
  xlab("id") +
  coord_flip()
```



<https://stackoverflow.com/a/9231857/15333580>

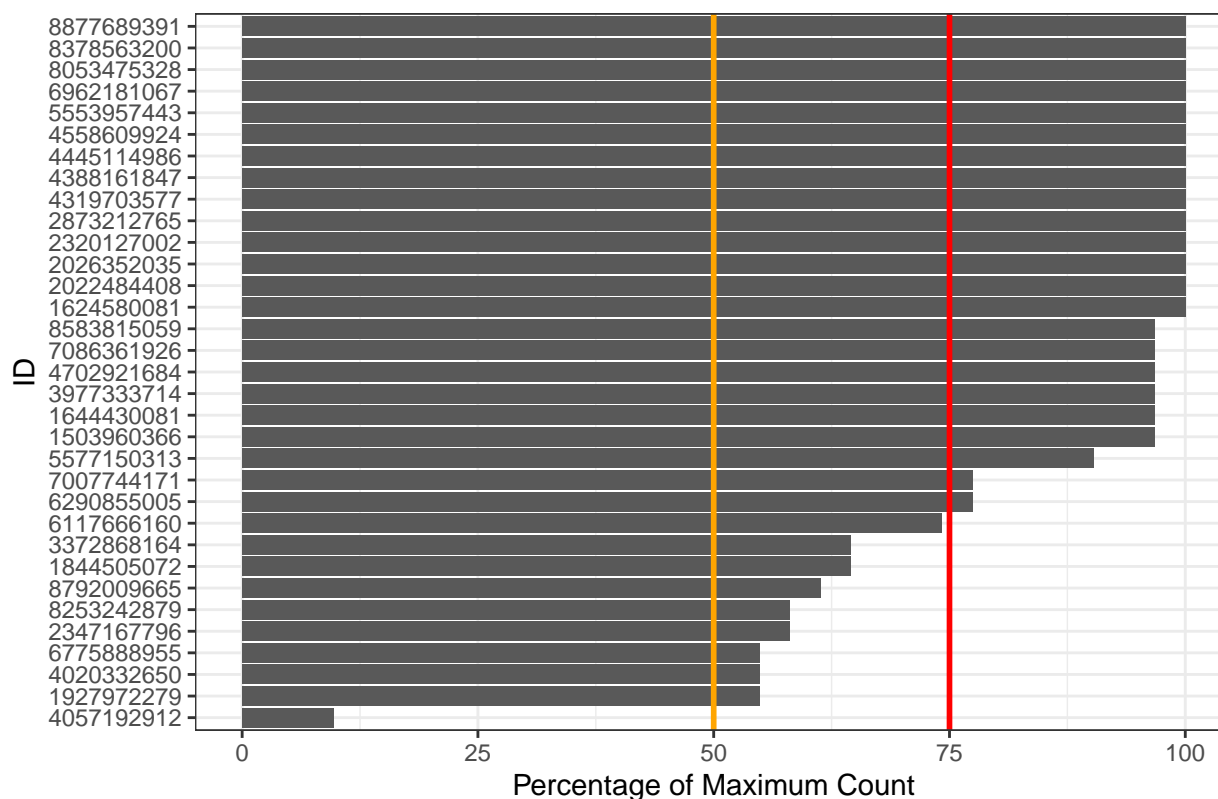
#reorder(id, id, length) takes the id variable, uses itself to determine the order, and uses the length

```
count_max_ratio <- daily_activity_clean %>%
  count(id) %>%
  rename(id = "id", count = "n") %>%
  mutate(percent_of_max = count / max(count) * 100) %>%
  arrange(desc(percent_of_max))
```

Create bar graph with percentage of entries compared to maximum

```
ggplot(count_max_ratio, aes(x = reorder(id, percent_of_max), y = percent_of_max)) +
  geom_bar(stat = "identity") +
  xlab("ID") +
  ylab("Percentage of Maximum Count") +
  ggtitle("Count by ID and Percentage of Maximum Count") +
  theme_bw() +
  theme(plot.title = element_text(hjust = 0.5)) +
  geom_hline(yintercept=50, color="orange", linewidth=1)+
  geom_hline(yintercept=75, color="red", linewidth=1)+
  coord_flip()
```

Count by ID and Percentage of Maximum Count



```
# percent_of_max > 75%
```

```
percent_of_max_top_75 <- filter(count_max_ratio, percent_of_max >=75)
percent_of_max_top_75
```

```
# A tibble: 23 x 3
```

	id	count	percent_of_max
	<chr>	<int>	<dbl>
1	1624580081	31	100
2	2022484408	31	100
3	2026352035	31	100
4	2320127002	31	100
5	2873212765	31	100
6	4319703577	31	100
7	4388161847	31	100
8	4445114986	31	100
9	4558609924	31	100
10	5553957443	31	100

```
# ... with 13 more rows
```

```
# percent_of_max < 75
```

```
percent_of_max_under_75 <- filter(count_max_ratio, percent_of_max < 75)
percent_of_max_under_75
```

```
# A tibble: 10 x 3
```

	id	count	percent_of_max
	<chr>	<int>	<dbl>


1	6117666160	23	74.2
2	1844505072	20	64.5
3	3372868164	20	64.5
4	8792009665	19	61.3
5	2347167796	18	58.1
6	8253242879	18	58.1
7	1927972279	17	54.8
8	4020332650	17	54.8
9	6775888955	17	54.8
10	4057192912	3	9.68

```
daily_activity_clean$activity_date %>% summary()
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
"2016-04-12"	"2016-04-18"	"2016-04-26"	"2016-04-26"	"2016-05-03"	"2016-05-12"

```
ggplot(data=daily_activity_clean , aes(x = activity_date)) +
  geom_histogram(binwidth = 1, color = "black", fill = "lightblue") +
  labs(x = "Activity Date", y = "Frequency", title = "Distribution of Activity Date")
```

Distribution of Activity Date




Observations:

- It appears that there is missing activity data towards the end of the available period, specifically in the beginning of May.


Investigate if the missing activity data coincides with the absence of entries for certain user IDs.

```
ggplot(data=subset(daily_activity_clean, id %in% percent_of_max_top_75$id), aes(x = activity_date)) +
  geom_histogram(binwidth = 1, color = "black", fill = "lightblue") +
```

```
labs(x = "Activity Date", y = "Frequency", title = "Distribution of Activity Date For IDs with Above 75% of Entries")
```



```
ggplot(data=subset(daily_activity_clean, id %in% percent_of_max_under_75$id), aes(x = activity_date)) +  
  geom_histogram(binwidth = 1, color = "black", fill = "lightblue") +  
  labs(x = "Activity Date", y = "Frequency", title = "Distribution of Activity Date For IDs with under 75% of Entries")
```



- Users with more than 75% of data consistently report activity dates, while those with less than 75% of data show a decline in reporting starting from the end of April. The decline in Activity Date seems to be primarily due to a lack of data reporting from some users during that period.


Bivariate analysis

```
corr <- cor(select_if(daily_activity_clean, is.numeric))

ggcorrplot(corr,
  hc.order = TRUE,
  type = "lower",
  lab = TRUE,
  colors = c("firebrick", "white", "royalblue"),
  lab_size = 4,
  lab_col = "black",
  title = "Correlation Between Numerical Variables")
```

Correlation between numerical variables

Correlation Between Numerical Variables



<https://rdr.io/github/microresearcher/MicroVis/man/ggcorrplot.html>

sedentary_minutes; sedentary_active_distance lightly_active_minutes; light_active_distance
fairly_active_minutes; moderately_active_distance very_active_minutes; very_active_distance

```
# Compute correlation matrix
```

```
corr_matrix <- corr
```

```
# Set the threshold for correlation
```

```
threshold <- 0.60
```

```
# Find pairs of highly correlated variables
```

```
high_cor_pairs <- which(abs(corr_matrix) > threshold & lower.tri(corr_matrix, diag = FALSE), arr.ind = TRUE)
```

```
# Extract the variable names and correlation coefficients for the correlated pairs
```

```
variable_names <- colnames(corr_matrix)
```

```
cor_values <- as.vector(corr_matrix[high_cor_pairs])
```

```
# Create a data frame to store the correlated pairs and their correlation coefficients
```

```
cor_data <- data.frame(
```

```

Variable1 = variable_names[high_cor_pairs[, 1]],
Variable2 = variable_names[high_cor_pairs[, 2]],
Correlation = cor_values
)

# Sort the correlated pairs by correlation coefficient in descending order
sorted_cor_data <- cor_data[order(-cor_data$Correlation), ]

# Remove the index
row.names(sorted_cor_data) <- NULL

# Display the sorted correlated variable pairs in the dataframe
print(sorted_cor_data)

```

	Variable1	Variable2	Correlation
1	tracker_distance	total_distance	0.9993982
2	total_distance	total_steps	0.9826464
3	tracker_distance	total_steps	0.9819287
4	fairly_active_minutes	moderately_active_distance	0.9448137
5	lightly_active_minutes	light_active_distance	0.8463101
6	very_active_minutes	very_active_distance	0.8215184
7	very_active_distance	total_distance	0.8088356
8	very_active_distance	tracker_distance	0.8087337
9	very_active_distance	total_steps	0.7544861
10	very_active_minutes	total_distance	0.6755673
11	very_active_minutes	tracker_distance	0.6751272
12	very_active_minutes	total_steps	0.6639646
13	calories	tracker_distance	0.6246510
14	calories	total_distance	0.6242380
15	calories	very_active_minutes	0.6122349
16	light_active_distance	total_steps	0.6048838

- Total_distance, tracker_distance, and total_steps are highly correlated, so we will retain only total distance and total_steps as they provide similar information.
- The following minute and distance types are correlated. Which indicates that they report different aspects of the same activity, this is time or distance:
 - lightly_active_minutes and light_active_distance (corr = 0.85)
 - fairly_active_minutes and moderately_active_distance (corr = 0.94)
 - very_active_minutes and very_active_distance (corr = 0.82)
- There is a moderately high correlation between the time spent during very active periods and the total number of steps/total distance:
 - The correlation between very_active_minutes and total_distance is 0.68
 - The correlation between very_active_minutes and total_steps is 0.66
- There is a moderate correlation of 0.61 between the total duration of very active minutes and the estimated daily calories consumed.
- There is a moderate correlation of 0.62 between the total distance covered and the estimated daily calories consumed.
- There is a moderate correlation coefficient of 0.60 between the distance covered during light activity (light_active_distance) and the total number of steps taken (total_steps).


```

# List of correlated variable pairs
correlated_pairs <- list(c("total_steps", "total_distance"),
                        c("lightly_active_minutes", "light_active_distance"),
                        c("fairly_active_minutes", "moderately_active_distance"),
                        c("very_active_minutes", "very_active_distance"),
                        c("very_active_minutes", "total_distance"),
                        c("very_active_minutes", "total_steps"),
                        c("very_active_minutes", "calories"),
                        c("total_distance", "calories"),
                        c("light_active_distance", "total_steps"))

# Loop over each pair and create scatter plot
for (pair in correlated_pairs) {
  var1 <- pair[1]
  var2 <- pair[2]


  # Calculate averages
  avg_var1 <- mean(daily_activity_clean[[var1]], na.rm = TRUE)
  avg_var2 <- mean(daily_activity_clean[[var2]], na.rm = TRUE)

  # Create scatter plot using ggplot2 with aes()
  print(ggplot(data = daily_activity_clean, aes(x = !!sym(var1), y = !!sym(var2))) +
        geom_point() +
        geom_vline(xintercept = avg_var1, linetype = "dashed", color = "red") +
        geom_hline(yintercept = avg_var2, linetype = "dashed", color = "blue") +
        xlab(var1) + ylab(var2) +
        ggtitle(paste("Scatter Plot with Average Reference Lines of", var1, "vs", var2)))
}


```

Scatterplots of selected highly correlated variables pairs (>0.60)


Scatter Plot with Average Reference Lines of total_steps vs total_distance




Scatter Plot with Average Reference Lines of lightly_active_minutes vs light_active_distance




Scatter Plot with Average Reference Lines of fairly_active_minutes vs moderately_active_distance




Scatter Plot with Average Reference Lines of very_active_minutes vs very_active_distance




Scatter Plot with Average Reference Lines of very_active_minutes vs total_c




Scatter Plot with Average Reference Lines of very_active_minutes vs total




Scatter Plot with Average Reference Lines of very_active_minutes vs calories



Scatter Plot with Average Reference Lines of total_distance vs calories



Scatter Plot with Average Reference Lines of light_active_distance vs total_steps



User Behavior for the daily activity dataset

```
# Create a boxplot for total_steps
boxplot(daily_activity_clean$total_steps,
        main = "Boxplot of Total Steps",
        ylab = "Total Steps")

# Calculate the median and standard deviation
median_value <- median(daily_activity_clean$total_steps)
std_dev <- round(sd(daily_activity_clean$total_steps),2)

# Identify outliers
outliers <- boxplot.stats(daily_activity_clean$total_steps)$out


# Count the number of outliers
num_outliers <- length(outliers)

# Create the legend label with median, standard deviation, and outlier count
legend_label <- paste("Median:", median_value,
                     "\nStandard Deviation:", std_dev,
                     "\nOutliers:", num_outliers)

# Add the legend with median, standard deviation, and outlier count
legend("topright", legend = legend_label, pch = "", col = "black", bty = "n", cex = 0.85)
```

Total steps: Total number of steps taken.

Boxplot of Total Steps



```
# Steps averages by IDs
steps_df <- daily_activity_clean %>%
  group_by(id) %>%
  summarise(average_steps = mean(total_steps), median_steps = median(total_steps), n = n())

steps_df
```

```
# A tibble: 33 x 4
  id       average_steps median_steps    n
  <chr>         <dbl>         <dbl> <int>
1 1503960366      12521.         12438    30
2 1624580081       5744.          4026    31
3 1644430081       7283.          6684.   30
4 1844505072       3999.          4036.   20
5 1927972279       1671.          1675    17
6 2022484408      11371.         11548    31
7 2026352035       5567.          5528    31
8 2320127002       4717.          5057    31
9 2347167796       9520.          9781    18
10 2873212765       7556.          7762    31
# ... with 23 more rows
```

```
# Calculate percentages for the average column
at_least_10k_avg <- sum(steps_df$average_steps >= 10000) / nrow(steps_df) * 100
between_5K_10K_avg <- sum(steps_df$average_steps >= 5000 & steps_df$average_steps < 10000) / nrow(steps_df) * 100
below_5k_avg <- sum(steps_df$average_steps < 5000) / nrow(steps_df) * 100
```

```
# Calculate percentages for the median column
at_least_10k_med <- sum(steps_df$median_steps >= 10000) / nrow(steps_df) * 100
between_5K_10K_med <- sum(steps_df$median_steps >= 5000 & steps_df$median_steps < 10000) / nrow(steps_df) * 100
below_5k_med <- sum(steps_df$median_steps < 5000) / nrow(steps_df) * 100
```

```
# Create a data frame for the steps categories
percentage_steps_df <- data.frame(
  Category = c("Below 5,000", "Between 5,000 and 10,000", "At least 10,000"),
```


```
Percentage_Average = round(c(below_5k_avg, between_5K_10K_avg, at_least_10k_avg)),
Percentage_Median = round(c(below_5k_med, between_5K_10K_med, at_least_10k_med)))
percentage_steps_df
```

	Category	Percentage_Average	Percentage_Median
1	Below 5,000	21	21
2	Between 5,000 and 10,000	58	52
3	At least 10,000	21	27

```
# Convert Category to a factor with custom factor levels
percentage_steps_df$Category <- factor(percentage_steps_df$Category, levels = c("Below 5,000", "Between
# Create a bar plot using ggplot
ggplot(percentage_steps_df, aes(x = Category, y = Percentage_Average)) +
  geom_bar(stat = "identity", fill = "blue") +
  labs(x = "Average Total Steps", y = "Percentage of Users", title = "58% of Users Average 5,000-10,000
  geom_text(aes(label = paste0(Percentage_Average, "%")), vjust = -0.5, color = "black") +
  ylim(0, 100) + theme_minimal() + theme(panel.grid = element_blank())
```

58% of Users Average 5,000–10,000 Step Daily

Only 21% Achieve the 10,000–Step Goal



```
# Create a boxplot for total_distance
boxplot(daily_activity_clean$total_distance,
        main = "Boxplot of Total Distance",
        ylab = "Total Distance")

# Calculate the median and standard deviation
median_value <- median(daily_activity_clean$total_distance)
```



```

std_dev <- sd(daily_activity_clean$total_distance)

# Identify outliers
outliers <- boxplot.stats(daily_activity_clean$total_distance)$out

# Count the number of outliers
num_outliers <- length(outliers)


# Create the legend label with median, standard deviation, and outlier count
legend_label <- paste("Median:", round(median_value, 2),
                      "\nStandard Deviation:", round(std_dev, 2),
                      "\nOutliers:", num_outliers)

# Add the legend with median, standard deviation, and outlier count
legend("topright", legend = legend_label, pch = "", col = "black", bty = "n")

```

Total Distance: Total kilometers tracked.

Boxplot of Total Distance



```

# Total distance by IDs
t_distance_df <- daily_activity_clean %>%
  group_by(id) %>%
  summarise(average_t_distance = mean(total_distance), median_t_distance = median(total_distance), n = n())

t_distance_df

```

A tibble: 33 x 4

	id	average_t_distance	median_t_distance	n
	<chr>	<dbl>	<dbl>	<int>
1	1503960366	8.07	8.08	30
2	1624580081	3.91	2.62	31
3	1644430081	5.30	4.86	30
4	1844505072	2.64	2.68	20
5	1927972279	1.16	1.16	17
6	2022484408	8.08	8.29	31
7	2026352035	3.45	3.45	31

```

8 2320127002          3.19          3.41    31
9 2347167796          6.36          6.54    18
10 2873212765          5.10          5.24    31
# ... with 23 more rows

```

```
# Calculate percentages for the average column
```

```
at_least_10_avg<- sum(t_distance_df$average_t_distance>= 10) / nrow(t_distance_df) * 100
```

```
between_5_10_avg <- sum(t_distance_df$average_t_distance >= 5 & t_distance_df$average_t_distance < 10) /
```

```
below_5_avg <- sum(t_distance_df$average_t_distance < 5) / nrow(t_distance_df) * 100
```

```
# Create a data frame for the distance categories
```

```
percentage_t_distance_df<- data.frame(
```

```
  Category = c("Below 5 km", "Between 5 and 10 km", "At least 10 km"),
```

```
  Percentage_Average = round(c(below_5_avg, between_5_10_avg , at_least_10_avg)))
```

```
percentage_t_distance_df
```

```

      Category Percentage_Average
1      Below 5 km              39
2 Between 5 and 10 km           55
3      At least 10 km              6

```

```
# Convert Category to a factor with custom factor levels
```

```
percentage_t_distance_df$Category <- factor(percentge_t_distance_df$Category, levels = c("Below 5 km",
```

```
# Create a bar plot using ggplot
```

```
ggplot(percentge_t_distance_df, aes(x = Category, y = Percentage_Average)) +
```

```
  geom_bar(stat = "identity", fill = "pink") +
```


```
  labs(x = "Average Total Distance", y = "Percentage of Users", title = "55% of Users Average 5-10 Kilometers")
```

```
  geom_text(aes(label = paste0(Percentage_Average, "%")), vjust = -0.5, color = "black") +
```

```
  ylim(0, 100) + theme_minimal() +theme(panel.grid = element_blank())
```

55% of Users Average 5–10 Kilometers Daily

10,000 steps is approximately equal to covering 5 miles (or 8 kilometers)



```
# Create a boxplot for sedentary_minutes
boxplot(daily_activity_clean$sedentary_minutes,
        main = "Boxplot of Sedentary Minutes",
        ylab = "Sedentary Minutes")

# Calculate the median and standard deviation
median_value <- median(daily_activity_clean$sedentary_minutes)
std_dev <- sd(daily_activity_clean$sedentary_minutes)

# Identify outliers
outliers <- boxplot.stats(daily_activity_clean$sedentary_minutes)$out


# Count the number of outliers
num_outliers <- length(outliers)

# Create the legend label with median, standard deviation, and outlier count
legend_label <- paste("Median:", round(median_value, 2),
                      "\nStandard Deviation:", round(std_dev, 2),
                      "\nOutliers:", num_outliers)

# Add the legend with median, standard deviation, and outlier count
legend("topright", legend = legend_label, pch = "", col = "black", bty = "n", cex = 0.80)
```

Sedentary Minutes: Total minutes spent in sedentary activity.

Boxplot of Sedentary Minutes



- These are high values for sedentary minutes. For instance, 1020 minutes is equivalent to 17 hours, and 1400 minutes is equivalent to 24 hours. After performing a quick search, it seems that the [Fitbit](#) uses 1400 as default for sedentary minutes when the device is not worn and it includes the sleeping time. SedentaryMinutes is total minutes spent in sedentary activity according to the data dictionary. See meta data section. Therefore, we need to subtract the times sleeping to obtain an more accurate estimate of daily sedentary minutes.

Sleep time is not considered sedentary time, so it was removed to determine the waking day and to allow the proportion of the day spent sedentary to be calculated

```
# Check sedentary_minutes stats
```

```
daily_activity_clean$sedentary_minutes %>% summary()
```

```
Min. 1st Qu. Median Mean 3rd Qu. Max.
 0.0   721.2 1020.5  955.2 1189.0 1440.0
```

```
outliers
```

```
[1] 2 13 0
```

```
# Count entries where sedentary minutes equal 1440
```

```
count_1440 <- sum(daily_activity_clean$sedentary_minutes == 1440)
```

```
# Output the count
```

```
count_1440
```

```
[1] 7
```

```
# Remove rows with sedentary minutes equal to the default value (1440) and outliers
```

```
daily_activity_clean <- filter(daily_activity_clean, !(sedentary_minutes %in% c(0, 2, 13, 1440)))
```

```
# Rename the column
```

```
daily_sleep_clean <- rename(daily_sleep_clean, activity_date = sleep_day)
```

```
# Join the datasets
```

```
joined_activity_sleep <- inner_join(daily_activity_clean, daily_sleep_clean, by = c("id", "activity_date"))
```

```

# Check missing values and duplicates
cat(
  "\n",
  "Missing values:",
  sum(is.na(joined_activity_sleep)),
  "\n",
  "Duplicate values:",
  sum(duplicated(joined_activity_sleep)),
  "\n",
  "Unique Ids:",
  n_distinct(joined_activity_sleep$id)
)

```

```

Missing values: 0
Duplicate values: 0
Unique Ids: 24

```

```

# Create a derived column for sedentary minutes that does not include sleep time
joined_activity_sleep <- joined_activity_sleep %>%
  mutate(
    sedentary_min_awake = sedentary_minutes - total_minutes_asleep,
    sedentary_hours_awake = sedentary_min_awake / 60,
    sedentary_percentage_diff = (sedentary_minutes - sedentary_min_awake) / sedentary_minutes * 100
  )

```

```

# Let us check the percentage difference of sedentary_minutes and the new column "sedentary_min_awake"

```

```

# Create a boxplot for sedentary_percentage_diff
boxplot(joined_activity_sleep$sedentary_percentage_diff,
  main = "Boxplot of Sedentary Percentage Difference",
  ylab = "Sedentary Percentage Difference")

# Calculate the median and standard deviation
median_value <- median(joined_activity_sleep$sedentary_percentage_diff)
std_dev <- sd(joined_activity_sleep$sedentary_percentage_diff)

# Identify outliers
outliers <- boxplot.stats(joined_activity_sleep$sedentary_percentage_diff)$out


# Count the number of outliers
num_outliers <- length(outliers)

# Create the legend label with median, standard deviation, and outlier count
legend_label <- paste("Median:", round(median_value, 2),
  "\nStandard Deviation:", round(std_dev, 2),
  "\nOutliers:", num_outliers)

# Add the legend with median, standard deviation, and outlier count
legend("topright", legend = legend_label, pch = "", col = "black", bty = "n", cex = 0.80)

```

Boxplot of Sedentary Percentage Difference



- The sedentary percentage difference has a median value of 59.95%, indicating a significant distinction between `sedentary_minutes` and `sedentary_min_away`. This suggests that the original column “`sedentary_minutes`” included the time asleep.

```
# Create a boxplot for sedentary_min_away
boxplot(joined_activity_sleep$sedentary_min_away,
        main = "Boxplot of Sedentary Minutes Awake",
        ylab = "Sedentary Minutes Awake")

# Calculate the median and standard deviation
median_value <- median(joined_activity_sleep$sedentary_min_away)
std_dev <- sd(joined_activity_sleep$sedentary_min_away)


# Identify outliers
outliers <- boxplot.stats(joined_activity_sleep$sedentary_min_away)$out

# Count the number of outliers
num_outliers <- length(outliers)

# Create the legend label with median, standard deviation, and outlier count
legend_label <- paste("Median:", round(median_value, 2),
                    "\nStandard Deviation:", round(std_dev, 2),
                    "\nOutliers:", num_outliers)

# Add the legend with median, standard deviation, and outlier count
legend("topright", legend = legend_label, pch = "", col = "black", bty = "n", cex = 0.80)
```

Boxplot of Sedentary Minutes Awake



- Observation: There appears to be an inconsistency in the data. The sedentary_minutes value is smaller than the total_minutes_asleep value, which is unexpected.

```
# Count the number of cases where sedentary_minutes is smaller than total_minutes_asleep
count <- sum(joined_activity_sleep$sedentary_minutes < joined_activity_sleep$total_minutes_asleep)
```

```
# Print the count
count
```

```
[1] 42
```

```
# Subset the dataset
```

```
subset_data <- joined_activity_sleep[joined_activity_sleep$sedentary_minutes < joined_activity_sleep$total_minutes_asleep, ]
```

```
# View the subsetted data
```

```
subset_data
```

```
# A tibble: 42 x 21
```

	id	activity~1	total~2	total~3	track~4	logge~5	very_~6	moder~7	light~8
	<chr>	<date>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	1503960366	2016-04-17	9705	6.48	6.48	0	3.19	0.780	2.51
2	1503960366	2016-05-08	10060	6.58	6.58	0	3.53	0.320	2.73
3	1644430081	2016-05-02	3758	2.73	2.73	0	0.0700	0.310	2.35
4	1844505072	2016-04-15	3844	2.54	2.54	0	0	0	2.54
5	1844505072	2016-04-30	4014	2.67	2.67	0	0	0	2.65
6	1844505072	2016-05-01	2573	1.70	1.70	0	0	0.260	1.45
7	1927972279	2016-04-12	678	0.470	0.470	0	0	0	0.470
8	2026352035	2016-04-23	12357	7.71	7.71	0	0	0	7.71
9	2026352035	2016-05-04	6564	4.07	4.07	0	0	0	4.07
10	2026352035	2016-05-06	8198	5.08	5.08	0	0	0	5.08

```
# ... with 32 more rows, 12 more variables: sedentary_active_distance <dbl>,
# very_active_minutes <dbl>, fairly_active_minutes <dbl>,
# lightly_active_minutes <dbl>, sedentary_minutes <dbl>, calories <dbl>,
# total_sleep_records <dbl>, total_minutes_asleep <dbl>,
# total_time_in_bed <dbl>, sedentary_min_away <dbl>,
# sedentary_hours_away <dbl>, sedentary_percentage_diff <dbl>, and
```

```

# abbreviated variable names 1: activity_date, 2: total_steps, ...
# Check column names of the subsetted data
subset_data %>%
select(sedentary_minutes, total_minutes_asleep, sedentary_min_away, calories,id, activity_date, total_

# A tibble: 42 x 9
  sedentary_~1 total~2 seden~3 calor~4 id activity~5 total~6 total~7 very~8
      <dbl>    <dbl>    <dbl>    <dbl> <chr> <date>      <dbl>    <dbl>    <dbl>
1         539        700       -161    1728 1503~ 2016-04-17    9705     6.48     38
2         574        594       -20    1740 1503~ 2016-05-08   10060     6.58     44
3         682        796      -114    2580 1644~ 2016-05-02    3758     2.73      1
4         527        644      -117    1725 1844~ 2016-04-15    3844     2.54      0
5         218        722     -504    1763 1844~ 2016-04-30    4014     2.67      0
6         585        590        -5    1541 1844~ 2016-05-01    2573     1.70      0
7         734        750       -16    2220 1927~ 2016-04-12      678     0.470     0
8         458        522       -64    1916 2026~ 2016-04-23   12357     7.71      0
9         530        538        -8    1658 2026~ 2016-05-04    6564     4.07      0
10        511        524       -13    1736 2026~ 2016-05-06    8198     5.08      0
# ... with 32 more rows, and abbreviated variable names 1: sedentary_minutes,
# 2: total_minutes_asleep, 3: sedentary_min_away, 4: calories,
# 5: activity_date, 6: total_steps, 7: total_distance, 8: very_active_minutes

dim(subset_data)

[1] 42 21

dim(joined_activity_sleep)

[1] 408 21

# Use anti_join() to return a new dataset that includes all rows from the first dataset except for the
clean_subset<- anti_join(joined_activity_sleep, subset_data)

Joining with `by = join_by(id, activity_date, total_steps, total_distance,
tracker_distance, logged_activities_distance, very_active_distance,
moderately_active_distance, light_active_distance, sedentary_active_distance,
very_active_minutes, fairly_active_minutes, lightly_active_minutes,
sedentary_minutes, calories, total_sleep_records, total_minutes_asleep,
total_time_in_bed, sedentary_min_away, sedentary_hours_away,
sedentary_percentage_diff)`

dim(clean_subset)

[1] 366 21

# Create a boxplot for sedentary_min_away
boxplot(clean_subset$sedentary_min_away,
        main = "Boxplot of Sedentary Minutes Awake",
        ylab = "Sedentary Minutes Awake")

# Calculate the median and standard deviation
median_value <- median(clean_subset$sedentary_min_away)
std_dev <- sd(clean_subset$sedentary_min_away)

# Identify outliers
outliers <- boxplot.stats(clean_subset$sedentary_min_away)$out

```



```


# Count the number of outliers
num_outliers <- length(outliers)

# Create the legend label with median, standard deviation, and outlier count
legend_label <- paste("Median:", round(median_value, 2),
  "\nStandard Deviation:", round(std_dev, 2),
  "\nOutliers:", num_outliers)

# Add the legend with median, standard deviation, and outlier count
legend("topright", legend = legend_label, pch = "", col = "black", bty = "n", cex = 0.80)

```

Boxplot of Sedentary Minutes Awake



Observation: By eliminating negative values from “sedentary_min_awake,” the resulting values now reflect a more realistic scenario.

```

# Total sedentary minutes awake by IDs
t_sedentary_df <- clean_subset %>%
  group_by(id) %>%
  summarise(average_sedentary_min_awake = mean(sedentary_min_awake),
    median_sedentary_min_awake = median(sedentary_min_awake), n = n())

t_sedentary_df

```

```

# A tibble: 23 x 4
  id          average_sedentary_min_awake median_sedentary_min_awake    n
  <chr>                <dbl>                <dbl> <int>
1 1503960366          442.                433     23
2 1644430081          873.                854      3
3 1927972279          704.                675      4
4 2026352035          181.                158     24
5 2320127002        1068                1068      1
6 2347167796          245.                220     13
7 3977333714          423.                420     28
8 4020332650          492.                440      8
9 4319703577          209.                148     23
10 4388161847          345                294     23

```

```

# ... with 13 more rows
dataset <- t_sedentary_df
column <- "average_sedentary_min_awake"
new_categories <- c("Below 200 minutes", "Between 200 and 400 minutes", "At least 400 minutes")

# Calculate percentages for the average column
below_200_avg <- sum(dataset[[column]] < 200) / nrow(dataset) * 100
between_200_400_avg <- sum(dataset[[column]] >= 200 & dataset[[column]] <= 400) / nrow(dataset) * 100
at_least_400_avg <- sum(dataset[[column]] >= 400) / nrow(dataset) * 100

# Create a data frame for the categories
percentage_sedentary_awake_df <- data.frame(
  Category = new_categories,
  Percentage_Average = round(c(below_200_avg, between_200_400_avg, at_least_400_avg))
)

# Convert Category to a factor with custom factor levels
percentage_sedentary_awake_df$Category <- factor(percentages_sedentary_awake_df$Category, levels = new_c


percentage_sedentary_awake_df

      Category Percentage_Average
1      Below 200 minutes           9
2 Between 200 and 400 minutes        43
3      At least 400 minutes        48

# Create a bar plot using ggplot
ggplot(percentages_sedentary_awake_df, aes(x = Category, y = Percentage_Average)) +
  geom_bar(stat = "identity", fill = "gray") +
  labs(x = "Average Total Sedentary Min Awake", y = "Percentage of Users",
       title = "48% of Users Have an Average of at Least 400 Daily Sedentary Minutes While Awake",
       subtitle = "200 Minutes are 3 hours and 20 minutes; 400 min are 6 hours and 40 min") +
  geom_text(aes(label = paste0(Percentage_Average, "%")), vjust = -0.5, color = "black") +
  ylim(0, 100) +
  theme_minimal() +
  theme(panel.grid = element_blank(), plot.title = element_text(size = 12), plot.subtitle = element_text

```

48% of Users Have an Average of at Least 400 Daily Sedentary Minutes While At Least 200 Minutes are 3 hours and 20 minutes; 400 min are 6 hours and 40 min



In a representative sample of U.S. adults, over two-thirds spent 6 + hours/day sitting, and more than half did not meet the recommended 150 min/week of physical activity. The study discovered that prolonged sitting for 6+ hours/day was associated with higher body fat percentages. While exceeding 150 min/week of physical activity was linked to lower body fat percentages, achieving recommended activity levels may not fully offset the increased body fat from prolonged sitting.

Jingwen Liao, Min Hu, Kellie Imm, Clifton J. Holmes, Jie Zhu, Chao Cao, Lin Yang. Association of daily sitting time and leisure-time physical activity with body fat among U.S. adults. Journal of Sport and Health Science, 2022. ISSN 2095-2546. <https://doi.org/10.1016/j.jshs.2022.10.001>. (<https://www.sciencedirect.com/science/article/pii/S2095254622001016>)

```
# Create a boxplot for calories
boxplot(daily_activity_clean$calories,
        main = "Boxplot of Calories",
        ylab = "Calories")

# Calculate the median and standard deviation
median_value <- median(daily_activity_clean$calories)
std_dev <- round(sd(daily_activity_clean$calories),2)

# Identify outliers
outliers <- boxplot.stats(daily_activity_clean$calories)$out


# Count the number of outliers
num_outliers <- length(outliers)
```

```
# Create the legend label with median, standard deviation, and outlier count
legend_label <- paste("Median:", median_value,
                      "\nStandard Deviation:", std_dev,
                      "\nOutliers:", num_outliers)

# Add the legend with median, standard deviation, and outlier count
legend("topright", legend = legend_label, pch = "", col = "black", bty = "n", cex = 0.85)
```

Calories: Total estimated energy expenditure (in kilocalories).

Boxplot of Calories



outliers

```
[1] 4552 4392 4501 4546 4900 4547 4398
```

```
# Calories averages by IDs
calories_df <- daily_activity_clean %>%
  group_by(id) %>%
  summarise(average_calories = mean(calories), median_calories = median(calories))
```

calories_df

A tibble: 33 x 3

	id	average_calories	median_calories
	<chr>	<dbl>	<dbl>
1	1503960366	1877.	1848
2	1624580081	1483.	1435
3	1644430081	2811.	2802.
4	1844505072	1732.	1752.
5	1927972279	2303.	2324
6	2022484408	2510.	2529
7	2026352035	1541.	1521
8	2320127002	1724.	1779
9	2347167796	2140.	2095
10	2873212765	1917.	1907

... with 23 more rows

```

# Calculate percentages for the average column
below_1600_avg <- sum(calories_df$average_calories < 1600) / nrow(calories_df) * 100
between_1600_2200_avg <- sum(calories_df$average_calories >= 1600 & calories_df$average_calories < 2200) / nrow(calories_df) * 100
between_2200_3000_avg <- sum(calories_df$average_calories >= 2200 & calories_df$average_calories < 3000) / nrow(calories_df) * 100
at_least_3000_avg <- sum(calories_df$average_calories >= 3000) / nrow(calories_df) * 100

# Calculate percentages for the median column
below_1600_med <- sum(calories_df$median_calories < 1600) / nrow(calories_df) * 100
between_1600_2200_med <- sum(calories_df$median_calories >= 1600 & calories_df$median_calories < 2200) / nrow(calories_df) * 100
between_2200_3000_med <- sum(calories_df$median_calories >= 2200 & calories_df$median_calories < 3000) / nrow(calories_df) * 100
at_least_3000_med <- sum(calories_df$median_calories >= 3000) / nrow(calories_df) * 100

# Create a data frame for the calories categories
percentage_calories_df <- data.frame(
  Category = c("Below 1,600", "Between 1,600 and 2,200", "Between 2,200 and 3,000", "At least 3,000"),
  Percentage_Average = round(c(below_1600_avg, between_1600_2200_avg, between_2200_3000_avg, at_least_3000_avg), 1),
  Percentage_Median = round(c(below_1600_med, between_1600_2200_med, between_2200_3000_med, at_least_3000_med), 1)
)

# Convert Category to a factor with custom factor levels
percentage_calories_df$Category <- factor(percentages_calories_df$Category, levels = c("Below 1,600", "Between 1,600 and 2,200", "Between 2,200 and 3,000", "At least 3,000"))

percentage_calories_df

```


	Category	Percentage_Average	Percentage_Median
1	Below 1,600	9	9
2	Between 1,600 and 2,200	42	36
3	Between 2,200 and 3,000	36	36
4	At least 3,000	12	18

```

# Create a bar plot using ggplot
ggplot(percentages_calories_df, aes(x = Category, y = Percentage_Average)) +
  geom_bar(stat = "identity", fill = "red") +
  labs(x = "Calorie Categories", y = "Percentage of Users",
       title = "42% of Users Have an Average Daily Calorie Expenditure Between 1,600 and 2,200.",
       subtitle = "Most females require 1,600 to 2,200 calories per day, as per the Dietary Guidelines") +
  geom_text(aes(label = paste0(Percentage_Average, "%")), vjust = -0.5, color = "black") +
  ylim(0, 100) +
  theme_minimal() +
  theme(panel.grid = element_blank(),
        plot.title = element_text(size = 12),
        plot.subtitle = element_text(size = 10))

```

42% of Users Have an Average Daily Calorie Expenditure Between 1,600 and 2,200
Most females require 1,600 to 2,200 calories per day, as per the Dietary Guidelines for Americans



“Females ages 19 through 30 require about 1,800 to 2,400 calories a day. Males in this age group have higher calorie needs of about 2,400 to 3,000 a day. Calorie needs for adults ages 31 through 59 are generally lower; most females require about 1,600 to 2,200 calories a day and males require about 2,200 to 3,000 calories a day.”

U.S. Department of Agriculture and U.S. Department of Health and Human Services. Dietary Guidelines for Americans, 2020-2025. 9th Edition. December 2020. Available at [DietaryGuidelines.gov/](https://www.dietaryguidelines.gov/)

Intensity Minutes: Time spent in one of four intensity categories.

- VeryActiveMinutes: Total minutes spent in very active activity
- FairlyActiveMinutes: Total minutes spent in moderate activity
- LightlyActiveMinutes: Total minutes spent in light activity
- SedentaryMinutes: Total minutes spent in sedentary activity

```
activity_minutes_df <- daily_activity_clean %>%
  group_by(id) %>%
  summarise(
    average_very_active_minutes = mean(very_active_minutes),
    average_fairly_active_minutes = mean(fairly_active_minutes),
    average_lightly_active_minutes = mean(lightly_active_minutes),
    average_sedentary_minutes = mean(sedentary_minutes)
  )

activity_minutes_df
```

A tibble: 33 x 5

```

  id          average_very_active_minutes average_fairly_activ~1 avera~2 avera~3
<chr>          <dbl>          <dbl>    <dbl>    <dbl>
1 1503960366      40          19.8      227.      828.
2 1624580081      8.68        5.81     153.     1258.
3 1644430081      9.57        21.4     178.     1162.
4 1844505072       0.2         2        179.     1115.
5 1927972279      2.41        1.41     70.4     1244.
6 2022484408     36.3        19.4     257.     1113.
7 2026352035      0.0968       0.258    257.      689.
8 2320127002      1.35         2.58     198.     1220.
9 2347167796     14.3        21.8     267.      727.
10 2873212765     14.1         6.13     308      1097.

```

```
# ... with 23 more rows, and abbreviated variable names
```

```
# 1: average_fairly_active_minutes, 2: average_lightly_active_minutes,
```

```
# 3: average_sedentary_minutes
```

```
# Define the custom order of legend items
```

```
custom_order <- c( "Very Active", "Fairly Active", "Lightly Active", "Sedentary")
```

```
# Create the stacked bar plot
```

```
ggplot(activity_minutes_df, aes(y = id)) +
```

```
  geom_bar(aes(x = average_sedentary_minutes, fill = "Sedentary"), stat = "identity", width = 0.5) +
```

```
  geom_bar(aes(x = average_lightly_active_minutes, fill = "Lightly Active"), stat = "identity", width = 0.5) +
```

```
  geom_bar(aes(x = average_fairly_active_minutes, fill = "Fairly Active"), stat = "identity", width = 0.5) +
```

```
  geom_bar(aes(x = average_very_active_minutes, fill = "Very Active"), stat = "identity", width = 0.5) +
```

```
  xlab("Minutes") +
```

```
  ylab("ID") +
```


```
  ggtitle("Average Activity Minutes by ID") +
```

```
  scale_fill_manual(name = "", values = c("Very Active" = "red", "Fairly Active" = "orange", "Lightly Active" = "yellow", "Sedentary" = "green")) +
```

```
  theme_minimal() +
```

```
  theme(legend.position = "bottom", panel.grid = element_blank())
```

Average Activity Minutes by ID



Calculate the average for each column

```
averages <- colMeans(activity_minutes_df[, c("average_very_active_minutes",
      "average_fairly_active_minutes",
      "average_lightly_active_minutes",
      "average_sedentary_minutes")])
```



```

# Calculate the total average
total_average <- sum(averages)

# Calculate the proportions
proportions <- averages / total_average

# Create the new dataframe with modified row names
overall_average_df <- data.frame(Average = averages,
                                Percentage = proportions * 100)

# Modify the row names
row_names <- c("Very Active", "Fairly Active", "Lightly Active", "Sedentary")
row.names(overall_average_df) <- row_names

# Print the new dataframe
overall_average_df


```

	Average	Percentage
Very Active	21.18923	1.744450
Fairly Active	14.29851	1.177156
Lightly Active	206.97366	17.039558
Sedentary	972.20431	80.038837

```

ggplot(overall_average_df, aes(x = Percentage, y = reorder(row.names(overall_average_df), Percentage),
                                stat = "identity", width = 0.7, show.legend = FALSE) +
  geom_text(aes(label = paste0(round(Percentage), "%")), hjust = -0.2, color = "black", size = 4) +
  ylab("Minutes Intensity") +
  xlab("Percentage") +
  ggtitle("Users' Overall Average Intensity Minutes Consist Primarily of Sedentary and Lightly Active Time") +
  scale_fill_manual(values = c("Very Active" = "red", "Fairly Active" = "orange", "Lightly Active" = "lightgreen", "Sedentary" = "lightblue")) +
  scale_x_continuous(labels = NULL) +
  theme_minimal() +
  theme(legend.position = "none", panel.grid = element_blank(), axis.text.y = element_text(size = 10))

```



“Analyzing each individual’s average calorie intake can provide insights into their individual dietary habits and patterns. By comparing the individual averages to the overall average, you can identify individuals who consume more or fewer calories compared to the group average. This comparison can help in understanding variations in calorie intake and potential factors influencing individual differences.”

```


# Define the custom order of legend items

custom_order <- c("Very Active", "Fairly Active", "Lightly Active", "Sedentary")

# Create the stacked horizontal bar chart
ggplot(overall_average_df, aes(x = Percentage, y = factor(1), fill = factor(row.names(overall_average_d
geom_bar(stat = "identity", width = 0.7) +
  xlab("Percentage") +
  ylab("") +
  ggtitle("Users' Overall Average Intensity Minutes Consist Primarily of Sedentary and Lightly Active T
  scale_fill_manual(
    name = "",
    values = c(
      "Very Active" = "red",
      "Fairly Active" = "orange",
      "Lightly Active" = "lightgreen",
      "Sedentary" = "lightblue"
    ),
    breaks = custom_order
  ) +
  guides(fill = guide_legend(reverse = TRUE)) + # Reverse the order of the legend
  theme_minimal() +
  theme(legend.position = "top",
        panel.grid = element_blank(),
        axis.text.y = element_blank(), # Remove the y-axis text
        plot.title = element_text(size = 12, margin = margin(b = 20))) + # Adjust the title size and m
  geom_vline(xintercept = 97, color = "black", linetype = "dashed") +
  annotate("text", x = 97, y = 1, label = " 97%", vjust = -5.5, hjust = 0.1)

```

Users' Overall Average Intensity Minutes Consist Primarily of Sedentary and Lightly Active



These indicators provide insights into activity levels, sedentary behavior, and calorie burn. They can help track progress, set goals, and evaluate user behavior over time. Remember to consider the specific context and goals of your analysis to select and customize the most relevant KPIs for your use case. The context I will use is the guidelines for physical activity and diet for Americans:

- U.S. Department of Health and Human Services. (2019). Physical Activity Guidelines

for Americans (2nd ed.). Available at https://health.gov/sites/default/files/2019-09/Physical_Activity_Guidelines_2nd_edition.pdf

- U.S. Department of Agriculture and U.S. Department of Health and Human Services. Dietary Guidelines for Americans, 2020-2025. 9th Edition. December 2020. Available at [DietaryGuidelines.gov/](https://www.dietaryguidelines.gov/)

EDA for daily_sleep_clean

```
str(daily_sleep_clean)
```

```
tibble [410 x 5] (S3: tbl_df/tbl/data.frame)
 $ id          : chr [1:410] "1503960366" "1503960366" "1503960366" "1503960366" ...
 $ activity_date : Date[1:410], format: "2016-04-12" "2016-04-13" ...
 $ total_sleep_records : num [1:410] 1 2 1 2 1 1 1 1 1 1 ...
 $ total_minutes_asleep: num [1:410] 327 384 412 340 700 304 360 325 361 430 ...
 $ total_time_in_bed   : num [1:410] 346 407 442 367 712 320 377 364 384 449 ...
```

- activity_date (sleep_day): Date on which the sleep event started.
- total_sleep_records: Number of recorded sleep periods for that day. Includes naps > 60 min.
- total_minutes_asleep: Total number of minutes classified as being “asleep”.
- total_time_in_bed: Total minutes spent in bed, including asleep, restless, and awake, that occurred during a defined sleep record.

```
#Sanity check: Verify that the value of total_time_in_bed is greater than total_minutes_asleep, as we w
daily_sleep_clean[daily_sleep_clean$total_time_in_bed < daily_sleep_clean$total_minutes_asleep,]
```

```
# A tibble: 0 x 5
# ... with 5 variables: id <chr>, activity_date <date>,
#   total_sleep_records <dbl>, total_minutes_asleep <dbl>,
#   total_time_in_bed <dbl>
```


Univariate analysis

```
numerical_cols <- daily_sleep_clean%>%
  select_if(is.numeric)


# plotting all numerical variables
col_names <- colnames(numerical_cols )
for (i in col_names) {
  suppressWarnings(print(
    ggplot(numerical_cols , aes(numerical_cols [[i]])) +
      geom_histogram(
        bins = 30,
        color = "black",
        fill = "gray",
        aes(y = ..density..)
      ) +
      geom_density(
        color = "blue",
        size = 1
      ) +
      xlab(i) + ylab("Count") +
      ggtitle(paste("Histogram with Density Plot of", i))
  ))
}
```

```
))  
}
```


Histogram with Density Plot of total_sleep_records



Histogram with Density Plot of total_minutes_asleep



Histogram with Density Plot of total_time_in_bed




Bivariate analysis

```
# Correlation between numerical variables
corr <- cor(select_if(daily_sleep_clean, is.numeric))


ggcorrplot(corr,
            hc.order = TRUE,
            type = "lower",
            lab = TRUE,
            colors = c("firebrick", "white", "royalblue"),
            lab_size = 4,
            lab_col = "black",
            title = "Correlation Between Numerical Variables")
```

Correlation between numerical variables



```
ggplot(data = daily_sleep_clean, aes(x = total_minutes_asleep, y = total_time_in_bed)) +
  geom_point()
```

Scatterplots of total_minutes_asleep vs total_time_in_bed




User Behavior for daily sleep dataset

```
frequency_table <- as.data.frame(table(daily_sleep_clean$total_sleep_records))
frequency_table$Percentage <- frequency_table$Freq / sum(frequency_table$Freq) * 100

ggplot(data = frequency_table, aes(x = Var1, y = Freq)) +
  geom_bar(stat = "identity", fill = "steelblue") +
  geom_text(aes(label = paste(Freq, " (", percent(Percentage / 100), "%)", sep = "")),
            hjust = 0.5, vjust = -0.4, color = "black") +
  labs(x = "Total Sleep Records", y = "Frequency",
       title = "Uncommon Napping: 89% of Sleep Records Indicate a Singular Sleep Period.",
       subtitle = "Includes naps > 60 min.") +
  theme_minimal() +
  theme(panel.grid = element_blank(),
        plot.title = element_text(size = 12),
        plot.subtitle = element_text(size = 10, margin = margin(b = 20)))
```

Uncommon Napping: 89% of Sleep Records Indicate a Singular Sleep Period.
Includes naps > 60 min.



```
# Create a boxplot for total_minutes_asleep
boxplot(daily_sleep_clean$total_minutes_asleep,
        main = "Boxplot of Total Minutes Asleep",
        ylab = "Total Minutes Asleep")

# Calculate the median and standard deviation
median_value <- median(daily_sleep_clean$total_minutes_asleep)
std_dev <- round(sd(daily_sleep_clean$total_minutes_asleep), 2)

# Identify outliers
outliers <- boxplot.stats(daily_sleep_clean$total_minutes_asleep)$out


# Count the number of outliers
num_outliers <- length(outliers)

# Create the legend label with median, standard deviation, and outlier count
legend_label <- paste("Median:", median_value,
                      "\nStandard Deviation:", std_dev,
                      "\nOutliers:", num_outliers)

# Add the legend with median, standard deviation, and outlier count
legend("topright", legend = legend_label, pch = "", col = "black", bty = "n", cex = 0.85)
```

Total minutes asleep

Boxplot of Total Minutes Asleep



```
# Sleep duration averages by IDs with standard deviation and count (n)
sleep_df <- daily_sleep_clean %>%
  group_by(id) %>%
  summarise(average_sleep_minutes = mean(total_minutes_asleep),
            standard_deviation_sleep_minutes = sd(total_minutes_asleep),
            n = n())
```

sleep_df

```
# A tibble: 24 x 4
  id          average_sleep_minutes standard_deviation_sleep_minutes    n
  <chr>                <dbl>                <dbl> <int>
1 1503960366          360.                100.     25
2 1644430081          294                335.      4
3 1844505072          652                66.4      3
4 1927972279          417                219.      5
5 2026352035          506.                42.3     28
6 2320127002           61                 NA        1
7 2347167796          447.                43.0     15
8 3977333714          294.                63.9     28
9 4020332650          349.                141.      8
10 4319703577          477.                114.     26
```

... with 14 more rows

```
# Drop ID "2320127002" due to insufficient data for computing mean and standard deviation.
```

```
sleep_df <- sleep_df %>%
  filter(id != "2320127002")
```

sleep_df

```
# A tibble: 23 x 4
  id          average_sleep_minutes standard_deviation_sleep_minutes    n
  <chr>                <dbl>                <dbl> <int>
1 1503960366          360.                100.     25
2 1644430081          294                335.      4
3 1844505072          652                66.4      3
```

```

4 1927972279          417          219.      5
5 2026352035          506.          42.3     28
6 2347167796          447.          43.0     15
7 3977333714          294.          63.9     28
8 4020332650          349.          141.      8
9 4319703577          477.          114.     26
10 4388161847          400.          146.     23

```

```
# ... with 13 more rows
```

```
# Calculate percentages for the average column
```

```
below_6_hours <- sum(sleep_df$average_sleep_minutes < 360) / nrow(sleep_df) * 100
```

```
between_6_7_hours <- sum(sleep_df$average_sleep_minutes >= 360 & sleep_df$average_sleep_minutes < 420) /
```

```
at_least_7_hours <- sum(sleep_df$average_sleep_minutes >= 420) / nrow(sleep_df) * 100
```

```
# Create a data frame for the sleep duration categories
```

```
percentage_sleep_df <- data.frame(
```

```
  Category = c("Below 6 hours", "Between 6 and 7 hours", "At least 7 hours"),
```

```
  Percentage_Average = round(c(below_6_hours, between_6_7_hours, at_least_7_hours))
```

```
)
```

```
# Convert Category to a factor with custom factor levels
```

```
percentage_sleep_df$Category <- factor(sleep_df$Category, levels = c("Below 6 hours", "Between 6 and 7 hours", "At least 7 hours"))
```

```
percentage_sleep_df
```

```

      Category Percentage_Average
1    Below 6 hours              30
2 Between 6 and 7 hours              22
3    At least 7 hours              48

```

```
str(sleep_df)
```

```
'data.frame':   3 obs. of  2 variables:
```

```
$ Category      : Factor w/ 3 levels "Below 6 hours",...: 1 2 3
```

```
$ Percentage_Average: num  30 22 48
```

```
ggplot(sleep_df, aes(x = Category, y = Percentage_Average)) +
```

```
  geom_bar(stat = "identity", fill = "purple") +
```

```
  labs(x = "Average Sleep Duration", y = "Percentage of Users",
```

```
        title = "52% of Users Get Less Than 7 Hours of Sleep on Average Daily") +
```


```
  geom_text(aes(label = paste0(Percentage_Average, "%")), vjust = -0.5, color = "black") +
```

```
  ylim(0, 100) +
```

```
  theme_minimal() +
```

```
  theme(panel.grid = element_blank(), plot.title = element_text(size = 12), plot.subtitle = element_text(size = 12))
```

52% of Users Get Less Than 7 Hours of Sleep on Average Daily



#Error bars

Convert average_sleep_minutes and standard_deviation_sleep_minutes to hours

```
sleep_df$average_sleep_hours <- sleep_df$average_sleep_minutes / 60
```

```
sleep_df$standard_deviation_sleep_hours <- sleep_df$standard_deviation_sleep_minutes / 60
```


Create a bar plot for each 'id' with error bars representing standard deviation

```
ggplot(sleep_df, aes(x = id, y = average_sleep_hours)) +  
  geom_bar(stat = "identity", fill = "skyblue", color = "black") +  
  geom_errorbar(aes(ymin = average_sleep_hours - standard_deviation_sleep_minutes / 60,  
                    ymax = average_sleep_hours + standard_deviation_sleep_minutes / 60),  
                width = 0.2, position = position_dodge(0.9), color = "black") +  
  labs(x = "ID", y = "Average Sleep Duration (hours)",  
        title = "Sleep Consistency: Average Sleep Duration with Error Bars",  
        subtitle = "Error bars represent the standard deviation around the mean.") +  
  theme_minimal() +  
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +  
  geom_hline(yintercept = 7, linetype = "dashed", color = "red") +  
  scale_y_continuous(breaks = seq(0, 12, 1)) # Adjust the range as needed
```

Sleep duration consistency

Sleep Consistency: Average Sleep Duration with Error Bars

Error bars represent the standard deviation around the mean.



```
# Calculate sleep duration averages and standard deviations in hours
sleep_df <- daily_sleep_clean %>%
  group_by(id) %>%
  summarise(n = n(),
            average_sleep_hours = mean(total_minutes_asleep) / 60,      # Convert minutes to hours
            average_time_in_bed_hours = mean(total_time_in_bed) / 60,
            standard_deviation_sleep_hours = sd(total_minutes_asleep) / 60,
            standard_deviation_time_in_bed_hours = sd(total_time_in_bed) / 60,
            ) %>%
  mutate(time_difference_hours = average_time_in_bed_hours - average_sleep_hours, # Calculate the time
         average_awake_in_bed_hours = time_difference_hours, # Rename column "awake_in_bed"
         sd_awake_in_bed_hours = sd(time_difference_hours)) # Calculate SD for "awake_in_bed" in hours
```

sleep_df

```
# A tibble: 24 x 9
  id          n average_sl~1 avera~2 stand~3 stand~4 time_~5 avera~6 sd_aw~7
  <chr>    <int>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
1 1503960366    25        6.00      6.39      1.67      1.63      0.382    0.382    1.08
2 1644430081     4         4.9       5.77      5.58      6.84      0.867    0.867    1.08
3 1844505072     3        10.9      16.0       1.11       0       5.15     5.15     1.08
4 1927972279     5         6.95      7.30      3.65      3.73      0.347    0.347    1.08
5 2026352035    28         8.44      8.96      0.704     0.707     0.524    0.524    1.08
6 2320127002     1         1.02      1.15      NA        NA        0.133    0.133    1.08
7 2347167796    15         7.45      8.19      0.716     0.827     0.742    0.742    1.08
8 3977333714    28         4.89      7.69      1.07      1.24      2.79     2.79     1.08
```

```

 9 4020332650      8      5.82   6.33   2.35   2.64   0.506   0.506   1.08
10 4319703577     26      7.94   8.37   1.90   1.97   0.422   0.422   1.08
# ... with 14 more rows, and abbreviated variable names 1: average_sleep_hours,
# 2: average_time_in_bed_hours, 3: standard_deviation_sleep_hours,
# 4: standard_deviation_time_in_bed_hours, 5: time_difference_hours,
# 6: average_awake_in_bed_hours, 7: sd_awake_in_bed_hours

# Drop ID "2320127002" due to insufficient data for computing mean and standard deviation.
sleep_df <- sleep_df %>%
  filter(id != "2320127002")
dim(sleep_df)

[1] 23  9

create_boxplots_in_one_output <- function(data_frame, columns_to_analyze, decimal_places = 2) {
  num_columns <- length(columns_to_analyze)
  num_rows <- ceiling(num_columns / 2)

  par(mfrow = c(num_rows, 2)) # Set the plotting layout

  for (i in 1:num_columns) {
    column_name <- columns_to_analyze[i]
    boxplot(data_frame[[column_name]],
            ylab = column_name)

    median_value <- median(data_frame[[column_name]])
    std_dev <- round(sd(data_frame[[column_name]]), decimal_places)
    outliers <- boxplot.stats(data_frame[[column_name]])$out
    num_outliers <- length(outliers)

    legend_label <- paste("Median:", round(median_value, decimal_places),
                          "\nSD:", std_dev,
                          "\nOutliers:", num_outliers)



    legend("topright", legend = legend_label, pch = "", col = "black", bty = "n", cex = 0.75)
  }

  par(mfrow = c(1, 1)) # Reset the plotting layout to default
}

# Columns to analyze
columns_to_analyze <- c("average_sleep_hours", "average_awake_in_bed_hours")



# Call the function to create boxplots in one output
create_boxplots_in_one_output(sleep_df, columns_to_analyze, decimal_places = 2)

```



```
# Columns to analyze
columns_to_analyze <- c("standard_deviation_sleep_hours", "sd_awake_in_bed_hours")

# Call the function
create_boxplots_in_one_output(sleep_df, columns_to_analyze, decimal_places = 2)
```



```
#Columns with outliers to remove
columns_with_outliers <- c("average_sleep_hours", "average_awake_in_bed_hours", "standard_deviation_sleep_hours")

# Function to remove outliers from a column
remove_outliers <- function(data, column_name) {
  outlier_bounds <- boxplot.stats(data[[column_name]])$out
  data_no_outliers <- data[!(data[[column_name]] %in% outlier_bounds), ]
  return(data_no_outliers)
}

# Loop through each column and remove outliers
for (col in columns_with_outliers) {
  sleep_df <- remove_outliers(sleep_df, col)
}
```

```
sleep_df
```

```
# A tibble: 17 x 9
```

	id	n	average_sl~1	avera~2	stand~3	stand~4	time_~5	avera~6	sd_aw~7
	<chr>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	1503960366	25	6.00	6.39	1.67	1.63	0.382	0.382	1.08
2	2026352035	28	8.44	8.96	0.704	0.707	0.524	0.524	1.08
3	2347167796	15	7.45	8.19	0.716	0.827	0.742	0.742	1.08
4	4020332650	8	5.82	6.33	2.35	2.64	0.506	0.506	1.08
5	4319703577	26	7.94	8.37	1.90	1.97	0.422	0.422	1.08
6	4388161847	23	6.67	7.05	2.43	2.58	0.384	0.384	1.08
7	4445114986	28	6.42	6.95	1.59	1.73	0.527	0.527	1.08
8	4702921684	27	6.96	7.30	1.11	1.13	0.346	0.346	1.08
9	5553957443	31	7.72	8.43	1.80	1.91	0.706	0.706	1.08
10	5577150313	26	7.2	7.68	1.65	1.79	0.477	0.477	1.08
11	6117666160	18	7.98	8.50	1.40	1.55	0.523	0.523	1.08
12	6775888955	3	5.83	6.15	1.68	1.60	0.322	0.322	1.08
13	6962181067	31	7.47	7.77	1.04	1.11	0.302	0.302	1.08
14	7086361926	24	7.55	7.77	1.16	1.18	0.222	0.222	1.08
15	8053475328	3	4.95	5.03	3.47	3.52	0.0778	0.0778	1.08
16	8378563200	31	7.42	8.10	1.28	1.45	0.680	0.680	1.08
17	8792009665	15	7.26	7.56	1.09	1.15	0.302	0.302	1.08

```
# ... with abbreviated variable names 1: average_sleep_hours,
```

```
# 2: average_time_in_bed_hours, 3: standard_deviation_sleep_hours,
```

```
# 4: standard_deviation_time_in_bed_hours, 5: time_difference_hours,
```




```
# 6: average_awake_in_bed_hours, 7: sd_awake_in_bed_hours
```

```
# Check if outliers were removed
```

```
columns_to_analyze <- c("average_sleep_hours", "average_awake_in_bed_hours", "standard_deviation_sleep_h
```

```
# Call the function to create boxplots in one output
```

```
create_boxplots_in_one_output(sleep_df, columns_to_analyze, decimal_places = 2)
```



Let us divide the users into irregular sleepers and regular sleepers. We will use the 75th percentile

Define the Threshold (e.g., using the 75th percentile)

```
threshold <- quantile(sleep_df$standard_deviation_sleep_hours, 0.75)
```

Create a new column "sleeper_type" based on the threshold

```
sleep_df$sleeper_type <- ifelse(sleep_df$standard_deviation_sleep_hours > threshold, "irregular", "regular")
```

```
sleep_df
```

A tibble: 17 x 10

	id	n	avera~1	avera~2	stand~3	stand~4	time_~5	avera~6	sd_aw~7	sleep~8
	<chr>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<chr>
1	150396~	25	6.00	6.39	1.67	1.63	0.382	0.382	1.08	regular
2	202635~	28	8.44	8.96	0.704	0.707	0.524	0.524	1.08	regular
3	234716~	15	7.45	8.19	0.716	0.827	0.742	0.742	1.08	regular
4	402033~	8	5.82	6.33	2.35	2.64	0.506	0.506	1.08	irregu~
5	431970~	26	7.94	8.37	1.90	1.97	0.422	0.422	1.08	irregu~
6	438816~	23	6.67	7.05	2.43	2.58	0.384	0.384	1.08	irregu~
7	444511~	28	6.42	6.95	1.59	1.73	0.527	0.527	1.08	regular
8	470292~	27	6.96	7.30	1.11	1.13	0.346	0.346	1.08	regular
9	555395~	31	7.72	8.43	1.80	1.91	0.706	0.706	1.08	regular
10	557715~	26	7.2	7.68	1.65	1.79	0.477	0.477	1.08	regular
11	611766~	18	7.98	8.50	1.40	1.55	0.523	0.523	1.08	regular
12	677588~	3	5.83	6.15	1.68	1.60	0.322	0.322	1.08	regular
13	696218~	31	7.47	7.77	1.04	1.11	0.302	0.302	1.08	regular
14	708636~	24	7.55	7.77	1.16	1.18	0.222	0.222	1.08	regular
15	805347~	3	4.95	5.03	3.47	3.52	0.0778	0.0778	1.08	irregu~
16	837856~	31	7.42	8.10	1.28	1.45	0.680	0.680	1.08	regular


```

17 879200~    15    7.26    7.56    1.09    1.15    0.302    0.302    1.08 regular
# ... with abbreviated variable names 1: average_sleep_hours,
# 2: average_time_in_bed_hours, 3: standard_deviation_sleep_hours,
# 4: standard_deviation_time_in_bed_hours, 5: time_difference_hours,
# 6: average_awake_in_bed_hours, 7: sd_awake_in_bed_hours, 8: sleeper_type

```

```

# sleep_type counts
table(sleep_df$sleeper_type)

```

```

irregular    regular
         4         13

```

```
sleep_df
```

```
# A tibble: 17 x 10
```

	id	n	avera~1	avera~2	stand~3	stand~4	time_~5	avera~6	sd_aw~7	sleeper~8
	<chr>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<chr>
1	150396~	25	6.00	6.39	1.67	1.63	0.382	0.382	1.08	regular
2	202635~	28	8.44	8.96	0.704	0.707	0.524	0.524	1.08	regular
3	234716~	15	7.45	8.19	0.716	0.827	0.742	0.742	1.08	regular
4	402033~	8	5.82	6.33	2.35	2.64	0.506	0.506	1.08	irregu~
5	431970~	26	7.94	8.37	1.90	1.97	0.422	0.422	1.08	irregu~
6	438816~	23	6.67	7.05	2.43	2.58	0.384	0.384	1.08	irregu~
7	444511~	28	6.42	6.95	1.59	1.73	0.527	0.527	1.08	regular
8	470292~	27	6.96	7.30	1.11	1.13	0.346	0.346	1.08	regular
9	555395~	31	7.72	8.43	1.80	1.91	0.706	0.706	1.08	regular
10	557715~	26	7.2	7.68	1.65	1.79	0.477	0.477	1.08	regular
11	611766~	18	7.98	8.50	1.40	1.55	0.523	0.523	1.08	regular
12	677588~	3	5.83	6.15	1.68	1.60	0.322	0.322	1.08	regular
13	696218~	31	7.47	7.77	1.04	1.11	0.302	0.302	1.08	regular
14	708636~	24	7.55	7.77	1.16	1.18	0.222	0.222	1.08	regular
15	805347~	3	4.95	5.03	3.47	3.52	0.0778	0.0778	1.08	irregu~
16	837856~	31	7.42	8.10	1.28	1.45	0.680	0.680	1.08	regular
17	879200~	15	7.26	7.56	1.09	1.15	0.302	0.302	1.08	regular

```

# ... with abbreviated variable names 1: average_sleep_hours,
# 2: average_time_in_bed_hours, 3: standard_deviation_sleep_hours,
# 4: standard_deviation_time_in_bed_hours, 5: time_difference_hours,
# 6: average_awake_in_bed_hours, 7: sd_awake_in_bed_hours, 8: sleeper_type

```

```
color_options <- c("#E69F00", "#0072B2") # Blue: "#0072B2", Orange: "#E69F00"
```

```
# Function to create the violin plot for a given y-axis column
```

```

create_violin_plot <- function(data, x_axis_col, y_axis_col) {
  ggplot(data, aes_string(x = x_axis_col, y = y_axis_col, fill = x_axis_col)) +
    geom_violin(scale = "width", draw_quantiles = c(0.25, 0.5, 0.75), trim = FALSE) +
    geom_boxplot(width = 0.1, fill = "white", color = "black") +
    labs(x = "Sleeper Type", y = y_axis_col, title = paste("Comparison", x_axis_col, "for", y_axis_col)) +
    scale_fill_manual(values = color_options) +
    theme_minimal()
}

```


```
# Call the function to create the violin plots for each column
```

```


for (col in c("average_sleep_hours", "standard_deviation_sleep_hours", "average_awake_in_bed_hours", "sd_
  plot <- create_violin_plot(sleep_df, "sleeper_type", col)
  print(plot)

```


```
}
```




Comparison sleeper_type for standard_deviation_sleep_hours



Comparison sleeper_type for average_awake_in_bed_hours





Observations:

- Regular sleepers tend to have higher median average sleep hours compared to irregular sleepers. This suggests that individuals classified as regular sleepers are likely getting more sleep on average than those categorized as irregular sleepers.
- Additionally, the spread of the “average_sleep_hours” for irregular sleepers appears to be wider, indicating more variability in their sleep duration. In contrast, the violin plot for regular sleepers shows a narrower spread, suggesting that their sleep duration is more consistent.
- Regular sleepers exhibit a slightly higher median average awake-in-bed duration compared to irregular sleepers.

Summary: Regular sleepers get more sleep on average, have a more consistent sleep duration, and slightly higher median awake-in-bed duration than irregular sleepers.

EDA minute_sleep_clean

```
str(minute_sleep_clean)
```

```
tibble [187,978 x 4] (S3: tbl_df/tbl/data.frame)
 $ id      : chr [1:187978] "1503960366" "1503960366" "1503960366" "1503960366" ...
 $ date    : POSIXct[1:187978], format: "2016-04-12 02:47:30" "2016-04-12 02:48:30" ...
 $ value   : Factor w/ 3 levels "1","2","3": 3 2 1 1 1 1 1 2 2 2 ...
 $ log_id  : num [1:187978] 1.14e+10 1.14e+10 1.14e+10 1.14e+10 1.14e+10 1.14e+10 ...
```

This data seems to come from the Classic Sleep Log (1 minute)

Value indicating the sleep state. 1 = asleep, 2 = restless, 3 = awake

For more detail check : [Fitbit data dictionary](#)

```
# Add labels to the value column
```

```
minute_sleep_clean$value <- factor(minute_sleep_clean$value, levels = c("1", "2", "3"), labels = c("asleep", "restless", "awake"))
```

```
minute_sleep_clean %>% summary()
```

id	date	value
Length:187978	Min. :2016-04-11 20:48:00.00	asleep :171960
Class :character	1st Qu.:2016-04-19 02:48:00.00	restless: 14002
Mode :character	Median :2016-04-26 21:48:00.00	awake : 2016
	Mean :2016-04-26 13:31:23.11	
	3rd Qu.:2016-05-03 23:47:00.00	
	Max. :2016-05-12 09:56:00.00	

```
log_id
```

Min. :1.137e+10
1st Qu.:1.144e+10
Median :1.150e+10
Mean :1.150e+10
3rd Qu.:1.155e+10
Max. :1.162e+10


```
# Assuming "value" column represents total sleep records
```

```
frequency_table <- as.data.frame(table(minute_sleep_clean$value))
```

```
frequency_table$Percentage <- frequency_table$Freq / sum(frequency_table$Freq) * 100
```

```
ggplot(data = frequency_table, aes(x = Var1, y = Freq)) +  
  geom_bar(stat = "identity", fill = "#008080") +  
  geom_text(aes(label = paste(Freq, " (", percent(Percentage / 100), "%)", sep = "")),  
            hjust = 0.5, vjust = -0.4, color = "black") +  
  labs(x = "Total Minutes Records", y = "Frequency",  
       title = "User Sleep States: 91% of Minutes Spent Asleep with Minimal Interruptions:",  
       subtitle = "Restlessness: 7.4% | Awake: 1.1%") +  
  theme_minimal() +  
  theme(panel.grid = element_blank(),  
        plot.title = element_text(size = 12),  
        plot.subtitle = element_text(size = 10, margin = margin(b = 20)))
```

User Sleep States: 91% of Minutes Spent Asleep with Minimal Interruptions:
Restlessness: 7.4% | Awake: 1.1%



EDA for hourly_activity_clean

```
str(hourly_activity_clean)
```

```
tibble [22,099 x 6] (S3: tbl_df/tbl/data.frame)
 $ id          : chr [1:22099] "1503960366" "1503960366" "1503960366" "1503960366" ...
 $ activity_hour : POSIXct[1:22099], format: "2016-04-12 00:00:00" "2016-04-12 01:00:00" ...
 $ calories      : num [1:22099] 81 61 59 47 48 48 48 47 68 141 ...
 $ total_intensity : num [1:22099] 20 8 7 0 0 0 0 0 13 30 ...
 $ average_intensity: num [1:22099] 0.333 0.133 0.117 0 0 ...
 $ step_total     : num [1:22099] 373 160 151 0 0 ...
```

- Calories integer Total number of estimated calories burned.
- TotalIntensity: integer Value calculated by adding all the minute-level intensity values that occurred within the hour.
- AverageIntensity: intensity state exhibited during that hour (TotalIntensity for that ActivityHour divided by 60)
- StepTotal: Total number of steps taken.

For more detail check : [Fitbit data dictionary](#)

```
hourly_df <- hourly_activity_clean
```

```
# Extract "am" or "pm" from the activity_hour column
hourly_df$am_pm <- ifelse(format(hourly_df$activity_hour, "%p") == "AM", "am", "pm")
```

```

#Add a column for the hour
hourly_df$hour <- format(hourly_df$activity_hour, "%H")

# Define colors for AM and PM
color_palette <- c("#FFA500", "#ADD8E6") # Orange for AM, Light Blue for PM

# Custom function to generate the boxplot with dynamically set y-axis limits
generate_boxplot <- function(data, y_var, y_label, limit_factor) {
  y_limit <- quantile(data[[y_var]], 0.95) * limit_factor

  ggplot(data, aes(x = hour, y = get(y_var), fill = am_pm)) +
    geom_boxplot(position = position_dodge(0.9), outlier.shape = NA) +
    scale_fill_manual(values = color_palette) +
    labs(title = paste("Median", y_label, "by Hour"),
         x = "Hour",
         y = paste("Median", y_label)) +
    guides(fill = guide_legend(title = NULL)) + # Remove legend title
    theme_minimal() +
    theme(panel.grid.major.x = element_blank()) +
    coord_cartesian(ylim = c(0, y_limit))
}


# Assuming your dataset is named 'hourly_df1'
data <- hourly_df

# Create the plots with dynamically adjusted y-axis limits
# Adjust the limit_factor as needed (e.g., 1.1, 1.2, etc.)
calories_plot <- generate_boxplot(data, "calories", "Calorie Burn", limit_factor = 1.4)
total_intensity_plot <- generate_boxplot(data, "total_intensity", "Total Intensity", limit_factor = 1.3)
step_total_avg_plot <- generate_boxplot(data, "step_total", "Total Steps", limit_factor = 1.3)


# Print the plots
print(calories_plot)

```

Median Calorie Burn by Hour




```
print(total_intensity_plot)
```

```
print(step_total_avg_plot)
```

Median Total Steps by Hour



Group by id, hour, and am_pm and summarize the columns

```
summary_data <- hourly_df %>%
  group_by(id, hour, am_pm) %>%
  summarize(
    calories_avg = mean(calories),
    calories_max = max(calories),
    calories_min = min(calories),
    total_intensity_avg = mean(total_intensity),
    total_intensity_max = max(total_intensity),
    total_intensity_min = min(total_intensity),
    average_intensity_avg = mean(average_intensity),
    step_total_avg = mean(step_total),
    observations_count = n()
  )
```

``summarise()`` has grouped output by 'id', 'hour'. You can override using the ``groups`` argument.

Define colors for AM and PM

```
color_palette <- c("#FFA500", "#ADD8E6") # Orange for AM, Light Blue for PM
```

Custom function to generate the bar plot with dynamically set y-axis limits

```
generate_bar_plot <- function(data, y_var, y_label, limit_factor) {
  y_limit <- max(data[[paste0(y_var, "_avg")]]) * limit_factor

  ggplot(data, aes(x = hour, y = get(paste0(y_var, "_avg")), fill = am_pm)) +
    geom_bar(stat = "identity", position = "dodge") +
    scale_fill_manual(values = color_palette) +
```

```

labs(title = paste("Average", y_label, "by Hour"),
     x = "Hour",
     y = paste("Average", y_label)) +
guides(fill = guide_legend(title = NULL)) + # Remove legend title
theme_minimal() +
theme(panel.grid.major.x = element_blank()) +
coord_cartesian(ylim = c(0, y_limit))
}


# Assuming your dataset is named 'summary_data'
data <- summary_data

# Create the bar plots with dynamically adjusted y-axis limits
calories_plot <- generate_bar_plot(data, "calories", "Calorie Burn", limit_factor = 1.2)
total_intensity_plot <- generate_bar_plot(data, "total_intensity", "Total Intensity", limit_factor = 1.2)
steps_plot <- generate_bar_plot(data, "step_total", "Steps Taken", limit_factor = 1.1)

# Print the plots
print(calories_plot)


```

Average Calorie Burn by Hour




```
print(total_intensity_plot)
```

Average Total Intensity by Hour



```
print(steps_plot)
```




```
# Function to create a box plot with customizable orientation and colors for a given y-axis column
create_custom_box_plot <- function(data, x_axis_col, y_axis_col, orientation = "ver", colors) {
  ggplot(data, aes_string(x = x_axis_col, y = y_axis_col, fill = x_axis_col)) +
    geom_boxplot(width = ifelse(orientation == "ver", 0.2, 0.5),
                 color = "black") + # Remove fill = "white"
  labs(x = "", y = y_axis_col,
       title = paste("Comparison", x_axis_col, "for", y_axis_col)) +
  scale_fill_manual(values = colors) +
  guides(fill = "none") + # Remove the legend for fill color
  theme_minimal() +
  if (orientation == "hor") coord_flip()
}


# Call the function with custom colors as a tuple and specify the x-axis label
custom_colors <- c("#FFA500", "#ADD8E6") # Orange for AM, Light Blue for PM

for (col in c("calories", "total_intensity", "step_total")) {
  plot <- create_custom_box_plot(hourly_df, x_axis_col = "am_pm", y_axis_col = col, orientation = "hor")
  print(plot)
}
```


Comparison am_pm for calories



Comparison am_pm for total_intensity



Comparison am_pm for step_total



With the code provided, you can gain several insights into user behavior and activity patterns based on the “hourly_df” dataset. Some of the insights you can obtain are:

Hourly Activity Patterns: You can observe how user activity varies throughout the day. The `calories_avg`, `total_intensity_avg`, and `average_intensity_avg` columns will give you the average calorie burn, total intensity, and average intensity for each hour, respectively. This can help identify peak activity hours and periods of lower activity.

Variability in Activity: The `calories_max`, `calories_min`, `total_intensity_max`, and `total_intensity_min` columns will provide information about the maximum and minimum values of calorie burn and total intensity recorded during each hour. This can help you understand the range of activity levels and how much the activity varies from hour to hour.

Average Steps per Hour: The `step_total_avg` column will give you the average number of steps taken during each hour. This can help you identify the typical step count during different times of the day.

Observations Count: The `observations_count` column will show the number of data points (observations) available for each “id,” “hour,” and “am_pm” group. This can help you assess the data density and identify hours with more or fewer data points, which may influence the reliability of the insights.

AM vs. PM Activity: The “am_pm” column indicates whether the activity occurred during the morning (AM) or afternoon/evening (PM). You can compare the activity patterns between these two periods and explore any differences in user behavior during these times.

Individual User Insights: By grouping the data by “id,” you can also gain insights into each individual user’s activity patterns. You can assess their average calorie burn, intensity, and steps during various hours.

Overall, these insights can help you understand how users engage in physical activity throughout the day, identify peak activity hours, and detect any patterns or trends in their behavior. This information can be valuable for designing personalized fitness plans, optimizing activity programs, and making data-driven

decisions to improve health and well-being.

EDA for seconds_heartrate_clean

```
str(seconds_heartrate_clean)
```

```
tibble [2,483,658 x 3] (S3: tbl_df/tbl/data.frame)
 $ id      : chr [1:2483658] "2022484408" "2022484408" "2022484408" "2022484408" ...
 $ date_time : POSIXct[1:2483658], format: "2016-04-12 07:21:00" "2016-04-12 07:21:05" ...
 $ heart_rate: num [1:2483658] 97 102 105 103 101 95 91 93 94 93 ...
```

```
seconds_heartrate_clean %>% summary()
```

	id	date_time	heart_rate
Length:	2483658	Min. : 2016-04-12 00:00:00.00	Min. : 36.00
Class :	character	1st Qu.: 2016-04-19 06:18:10.00	1st Qu.: 63.00
Mode :	character	Median : 2016-04-26 20:28:50.00	Median : 73.00
		Mean : 2016-04-26 19:43:52.24	Mean : 77.33
		3rd Qu.: 2016-05-04 08:00:20.00	3rd Qu.: 88.00
		Max. : 2016-05-12 16:20:00.00	Max. : 203.00

```
n_distinct(seconds_heartrate_clean$id)
```

```
[1] 14
```

```
# Group by 'id' and calculate the average heart rate for each user
```

```
average_heart_rate <- seconds_heartrate_clean %>%
```

```
  group_by(id) %>%
```

```
  summarise(average_heart_rate = mean(heart_rate, na.rm = TRUE))
```

```
# Bar Plot with smaller bars and custom breaks on the heart rate axis
```

```
bar_plot <- ggplot(average_heart_rate, aes(x = id, y = average_heart_rate)) +
```

```
  geom_bar(stat = "identity", fill = "#ADD8E6", width = 0.8) + # Adjust the width here (e.g., 0.5 for
```

```
  labs(x = "User ID", y = "Average Heart Rate", title = "Average Heart Rate for Each User") +
```

```
  theme_minimal() + coord_flip() +
```


```
  scale_y_continuous(breaks = seq(0, 100, 10), limits = c(0, 100)) + # Set custom breaks and limits for
```

```
  geom_hline(yintercept = 60, linetype = "dashed", color = "green") +
```

```
  geom_hline(yintercept = 100, linetype = "dashed", color = "green")
```

```
# Display the bar plot
```

```
print(bar_plot)
```

EDA for weight_logs_clean

```
str(weight_logs_clean)
```

```
tibble [67 x 8] (S3: tbl_df/tbl/data.frame)
 $ id          : chr [1:67] "1503960366" "1503960366" "1927972279" "2873212765" ...
 $ date_time   : POSIXct[1:67], format: "2016-05-02 23:59:59" "2016-05-03 23:59:59" ...
 $ weight_kg   : num [1:67] 52.6 52.6 133.5 56.7 57.3 ...
 $ weight_pounds : num [1:67] 116 116 294 125 126 ...
 $ fat         : num [1:67] 22 0 0 0 0 25 0 0 0 0 ...
 $ bmi         : num [1:67] 22.6 22.6 47.5 21.5 21.7 ...
 $ is_manual_report: logi [1:67] TRUE TRUE FALSE TRUE TRUE TRUE ...
 $ log_id      : num [1:67] 1.46e+12 1.46e+12 1.46e+12 1.46e+12 1.46e+12 ...
```

- Fat: Body fat percentage recorded.
- BMI: Measure of body mass index based on the height and weight in the participant's Fitbit.com profile.
- isManualReport: If the data for this weigh-in was done manually (TRUE), or if data was measured and synced directly to Fitbit.com from a connected scale (FALSE).

For more detail check : [Fitbit data dictionary](#)

```
weight_logs_clean %>% summary()
```

id	date_time	weight_kg
Length:67	Min. :2016-04-12 06:47:11.00	Min. : 52.60
Class :character	1st Qu.:2016-04-19 15:19:45.00	1st Qu.: 61.40
Mode :character	Median :2016-04-27 23:59:59.00	Median : 62.50

```

Mean      :2016-04-27 15:39:54.27   Mean      : 72.04
3rd Qu.   :2016-05-04 15:24:10.50   3rd Qu.   : 85.05
Max.      :2016-05-12 23:59:59.00   Max.      :133.50

weight_pounds    fat          bmi    is_manual_report
Min.      :116.0    Min.      : 0.0000    Min.      :21.45    Mode :logical
1st Qu.    :135.4    1st Qu.    : 0.0000    1st Qu.    :23.96    FALSE:26
Median     :137.8    Median     : 0.0000    Median     :24.39    TRUE :41
Mean       :158.8    Mean       : 0.7015    Mean       :25.19
3rd Qu.    :187.5    3rd Qu.    : 0.0000    3rd Qu.    :25.56
Max.       :294.3    Max.       :25.0000    Max.       :47.54

log_id
Min.      :1.460e+12
1st Qu.   :1.461e+12
Median    :1.462e+12
Mean      :1.462e+12
3rd Qu.   :1.462e+12
Max.      :1.463e+12


```

```
# Create boxplots for "bmi" and "weight_pounds"
```

```
columns_to_analyze <- c("bmi", "weight_pounds")
```

```
# Call the function to create boxplots
```

```
create_boxplots_in_one_output(weight_logs_clean, columns_to_analyze, decimal_places = 2)
```



```

entry_count <- weight_logs_clean %>%
  group_by(id, is_manual_report) %>%
  summarize(entry_count = n(), .groups = "keep") %>%
  arrange (- entry_count)

```

```
print(entry_count)
```

```
# A tibble: 8 x 3
```

```
# Groups:   id, is_manual_report [8]
```

```

  id          is_manual_report entry_count
<chr>        <lgl>             <int>
1 6962181067 TRUE                30
2 8877689391 FALSE                24

```

```

3 4558609924 TRUE 5
4 1503960366 TRUE 2
5 2873212765 TRUE 2
6 4319703577 TRUE 2
7 1927972279 FALSE 1
8 5577150313 FALSE 1

```

```

# Check users that reported fat percentage
weight_logs_clean %>% filter(fat != 0)

```

```

# A tibble: 2 x 8
  id      date_time      weight_kg weight_~1  fat  bmi is_ma~2 log_id
<chr>    <dtm>          <dbl>    <dbl> <dbl> <dbl> <lgl>    <dbl>
1 1503960366 2016-05-02 23:59:59    52.6    116.   22  22.6 TRUE    1.46e12
2 4319703577 2016-04-17 23:59:59    72.4    160.   25  27.5 TRUE    1.46e12
# ... with abbreviated variable names 1: weight_pounds, 2: is_manual_report

```

Observation: - Only two users reported fat percentage

```

# Calculate total entries
total_entries <- sum(entry_count$entry_count)

```

```

average_bmi_weight <- weight_logs_clean %>%
  group_by(is_manual_report) %>%
  summarize(mean_bmi = mean(bmi, na.rm = TRUE),
            mean_weight_pounds = mean(weight_pounds, na.rm = TRUE),
            entry_count = n(),
            entry_count_percentage = round((n()/total_entries)*100,2),
            .groups = "keep")

print(average_bmi_weight)

```

```

# A tibble: 2 x 5
# Groups:   is_manual_report [2]
  is_manual_report mean_bmi mean_weight_pounds entry_count entry_count_percent~1
<lgl>            <dbl>          <dbl>          <int>          <dbl>
1 FALSE           26.4           192.           26           38.8
2 TRUE            24.4           138.           41           61.2
# ... with abbreviated variable name 1: entry_count_percentage

```


Observation: - 61% of users log their weight manually, while 39% sync their weight from other devices

```

create_custom_box_plot(weight_logs_clean , x_axis_col = "is_manual_report", y_axis_col = "weight_pounds

```

Comparison is_manual_report for weight_pounds



```
# Use the remove_outliers function we created previously to remove the outliers
```

```
columns_with_outliers <- c("weight_pounds", "bmi")
```


```
# Loop through each column and remove outliers
```

```
for (col in columns_with_outliers) {  
  weight_logs_clean <- remove_outliers(weight_logs_clean, col)  
}
```

```
dim(weight_logs_clean)
```

```
[1] 64 8
```

```
create_custom_box_plot(weight_logs_clean , x_axis_col = "is_manual_report", y_axis_col = "weight_pounds")
```



Observations:

- It appears that users who log their weight data manually have a lower median weight than users who sync their weight from other devices.
- The previous observation should be viewed as exploratory and could benefit from additional data. The weight log dataset only has 68 entries; more data would be needed to evaluate these hunches.
- The lack of completeness in the weight log dataset could indicate a lack of user engagement.

next

- add percentages of entries
- Revise notebook -Check for grammar
- Organize references and check headings and formatting
- Write final recommendations __ Write limitations
- Work on report and final recommendations
- Check knit and convert to jupyternotebook
-

Reference:

- EDA: <https://rpubs.com/jovial/r>
- Histograms: <https://statisticsbyjim.com/basics/histograms/> <https://blog.minitab.com/en/3-things-a-histogram-can-tell-you>

Categorical, ordinal, interval, and ratio variables : https://www.graphpad.com/guides/prism/latest/statistics/the_different_kinds_of_variabl.htm

Add density line to histogram: <https://r-coder.com/density-plot-r>

- Error bars vs CI: <https://blogs.sas.com/content/iml/2019/10/09/statistic-error-bars-mean.html>
1. Plotting histograms with ggplot2:1.0,

Insights and recommendations

<https://www.cdc.gov/mmwr/volumes/68/wr/mm6823a1.htm>

<https://stackoverflow.com/questions/13035834/plot-every-column-in-a-data-frame-as-a-histogram-on-one-page-using-ggplot>

<https://stackoverflow.com/questions/13035834/plot-every-column-in-a-data-frame-as-a-histogram-on-one-page-using-ggplot>

Another source:

<https://www.kaggle.com/datasets/arashnic/fitbit/discussion/313589?search=data>

#paper <https://dl.acm.org/doi/pdf/10.1145/3339825.3394926>

this is it: physical inactivity. Plot a barplot with percentages. <https://www.cdc.gov/physicalactivity/data/inactivity-prevalence-maps/index.html#Race-Ethnicity>

information about physical activity guidelines (sex and age):

<https://www.cdc.gov/nchs/products/databriefs/db443.htm> Elgaddal N, Kramarow EA, Reuben C. Physical activity among adults aged 18 and over: United States, 2020. NCHS Data Brief, no 443. Hyattsville, MD: National Center for Health Statistics. 2022. DOI: <https://dx.doi.org/10.15620/cdc:120213>

————— showing notebook in github

convert to jupyter notebook option: <https://medium.datadriveninvestor.com/transforming-your-rmd-to-ipynb-file-r-markdown-to-python-jupyter-b1306646f50b>

Hey all,

Sorry if I'm misunderstanding here, but I have been knitting the .Rmd notebook to a .md file within RStudio, and it seems to display very well in GitHub. You can see an example in my repo to see if I'm on track with this thread.

The links below give the explanation. Short Version:

change “output=html_document” to “output=github_document”

knit the document push the .md file to GitHub instead of the .Rmd be sure to push the ‘_files’ folder to include any images https://rmarkdown.rstudio.com/github_document_format.html <https://gist.github.com/JoshuaTPierce/b919168421b40e06481080eb53c3fb2f>